

Introduction to Data Management

♪ *Join* together... right now... over me ♪

Shana Hutchison

Paul G. Allen School of Computer Science and Engineering
University of Washington, Seattle

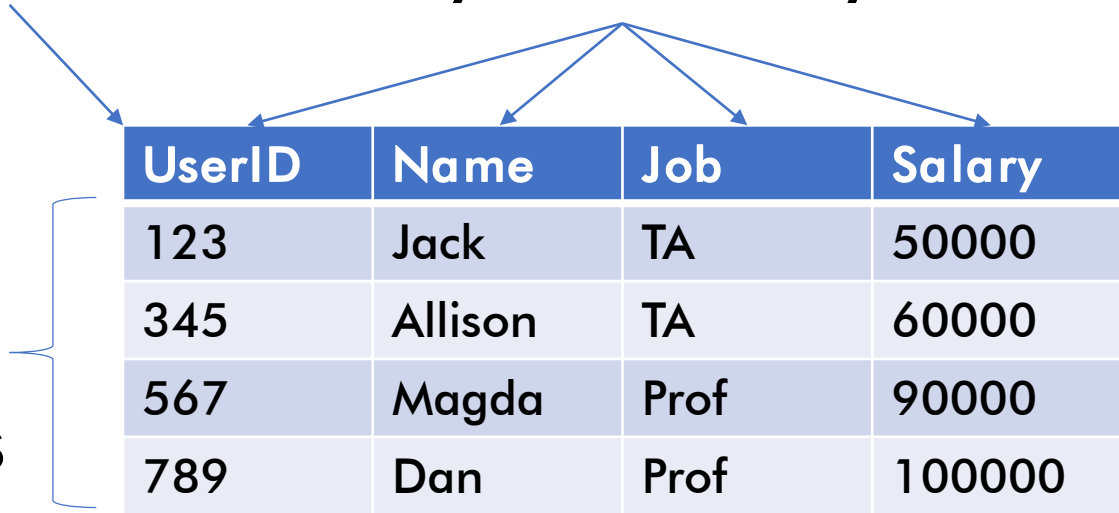
Recap: The Relational Model

- Flat tables, static and typed attributes, etc.
 - “It’s a spreadsheet with rules”

**Table/
Relation**

Columns/Attributes/Fields

**Rows/
Tuples/
Records**



The diagram illustrates the components of a table. A blue arrow points from the label 'Table/Relation' to the entire table structure. Another blue arrow points from the label 'Columns/Attributes/Fields' to the header row of the table. A third blue arrow points from the label 'Rows/Tuples/Records' to the first column of the table. The table itself has a blue header row and four data rows.

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Recap: FROM-WHERE-SELECT

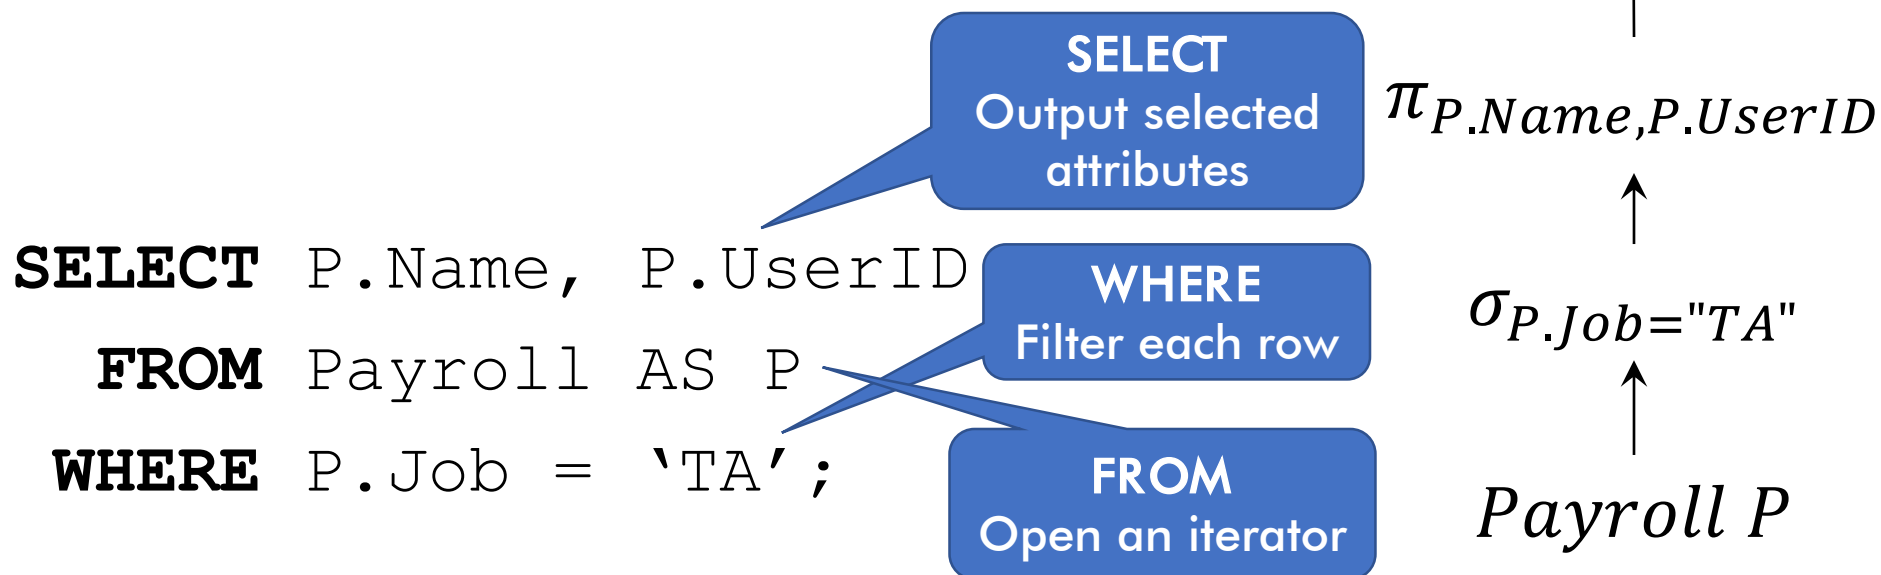
Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)



RA: FROM-WHERE-SELECT

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

for each row in P:

if (row.Job == 'TA'):

output (row.Name, row.UserID)

1-arg Projection op (Greek "pi")
Project input onto list of attributes

1-arg Selection op (Greek "sigma")
Filter rows of input with condition

0-arg Input op
Produce an input relation

SELECT
Output selected
attributes

WHERE
Filter each row

FROM
Open an iterator

(Output)
↑
 $\pi_{P.Name, P.UserID}$
↑
 $\sigma_{P.Job="TA"}$
↑
Payroll P

WHERE Practice

```
SELECT *  
  FROM Payroll AS P  
 WHERE P.UserID > 200;
```

UserID	Name	Job	Salary
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

$\sigma_{UserID > 200}$

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

WHERE Practice

UserID	Name	Job	Salary
123	Jack	TA	50000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT *  
  FROM Payroll AS P  
 WHERE P.UserID < 200  
        OR P.Job = "Prof";
```

$\sigma_{UserID > 200 \text{ OR Job} = \text{"Prof"}}$

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

SELECT Practice

```
SELECT P.Name, P.Salary/2
FROM Payroll AS P;
```

Name	Salary
Jack	25000
Allison	30000
Magda	45000
Dan	50000

$\pi_{Name, Salary/2}$

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

SELECT Practice

Duplicates!
May occur in output of
an operator.
Never in an input table.

```
SELECT P.Job  
FROM Payroll AS P;
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

π_{Job}


Job
TA
TA
Prof
Prof

DISTINCT

1-arg Duplicate Elim op
(Greek "delta")

```
SELECT DISTINCT P.Job  
FROM Payroll AS P;
```

Next lecture:
DISTINCT is equivalent to
GROUP BY with all attributes



π_{Job}

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

ORDER BY

1-arg Sort op
(Greek "tau")

```
SELECT *  
FROM Payroll AS P  
ORDER BY Job, Name;
```

Use DESC for
reverse order

UserID	Name	Job	Salary
789	Dan	Prof	100000
567	Magda	Prof	90000
345	Allison	TA	60000
123	Jack	TA	50000

$\tau_{Job, Name}$

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

CREATE TABLE

```
CREATE TABLE Payroll (  
    UserID int, Name text, Job text, Salary int);  
INSERT INTO Payroll VALUES  
    (123, 'Jack', 'TA', 50000),  
    (123, 'Allison', 'TA', 60000),  
    (123, 'Magda', 'Prof', 90000),  
    (123, 'Dan', 'Prof', 100000);
```

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

- FROM-WHERE-SELECT
- DISTINCT, ORDER BY, CREATE TABLE

Now you can do hw1! Next up:

- Keys
- Foreign Keys
- Joins + RA

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.

Definitely not a key

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.

Good candidate
for a key

Definitely not a key

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.

Is this a good candidate for a key?

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.



Is this a good candidate for a key?

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.



Is this a good candidate for a key?

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
913	Peter	TA	60000

Keys

Key

A **Key** is one or more attributes that uniquely identify a row.

Data comes from the real world
so models ought to reflect that

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
913	Peter	TA	60000

```
CREATE TABLE Payroll (  
    UserID INT,  
    Name TEXT,  
    Job TEXT,  
    Salary INT);
```

Payroll(UserId, Name, Job, Salary)

Unique Identifier

```
CREATE TABLE Payroll (  
  UserID INT,  
  Name TEXT,  
  Job TEXT,  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

Unique Identifier

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name TEXT,  
  Job TEXT,  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

Keys

```
CREATE TABLE Payroll (  
    UserID INT,  
    { Name TEXT,  
      Job TEXT,  
      Salary INT);
```

Suppose the set
of attributes
{Name, Job,
Salary} is a key

Payroll(UserId, Name, Job, Salary)

Keys

```
CREATE TABLE Payroll (  
    UserID INT,  
    { Name TEXT,  
      Job TEXT,  
      Salary INT,  
    PRIMARY KEY (Name, Job, Salary) );
```

Suppose the set
of attributes
{Name, Job,
Salary} is a key

Payroll(UserId, Name, Job, Salary)

Foreign Keys

- Databases can hold multiple tables
- How to capture relationships *between* tables?

Payroll

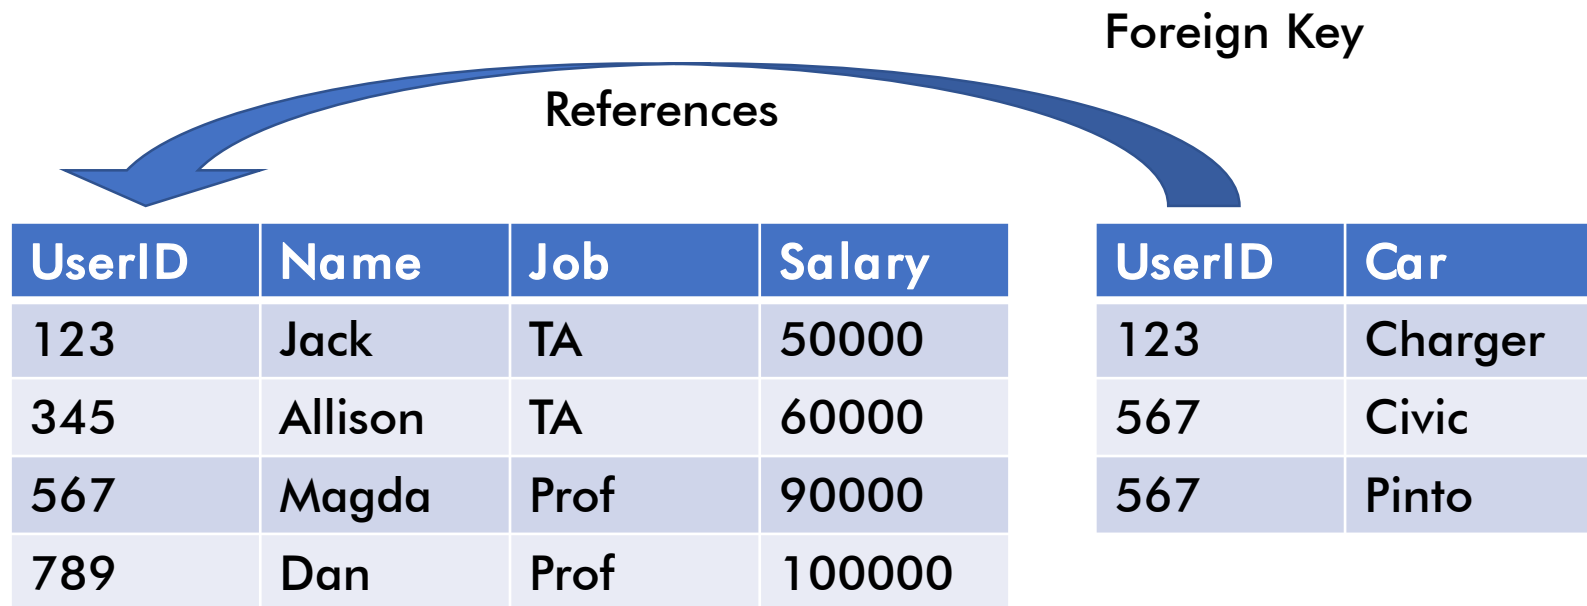
UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Foreign Keys

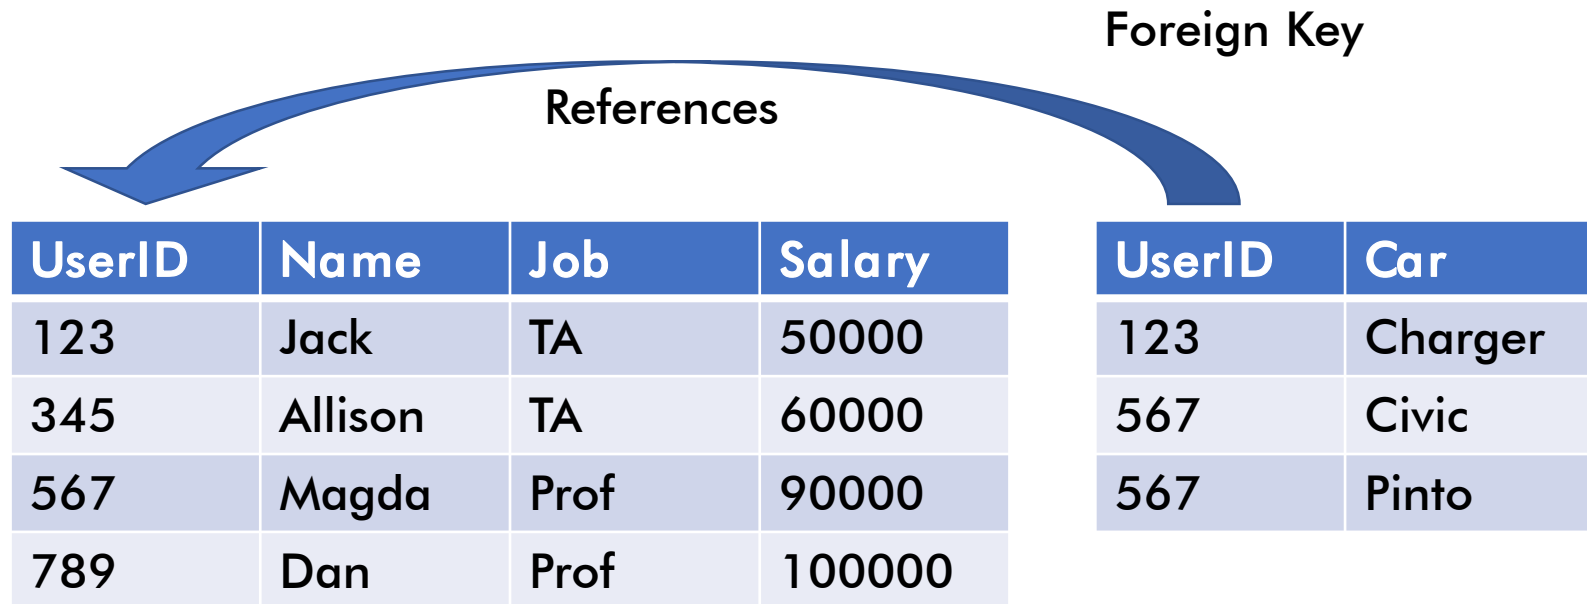
- Databases can hold multiple tables
- How to capture relationships *between* tables?



Foreign Keys

Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



Foreign Keys

Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.

Is this valid?

References

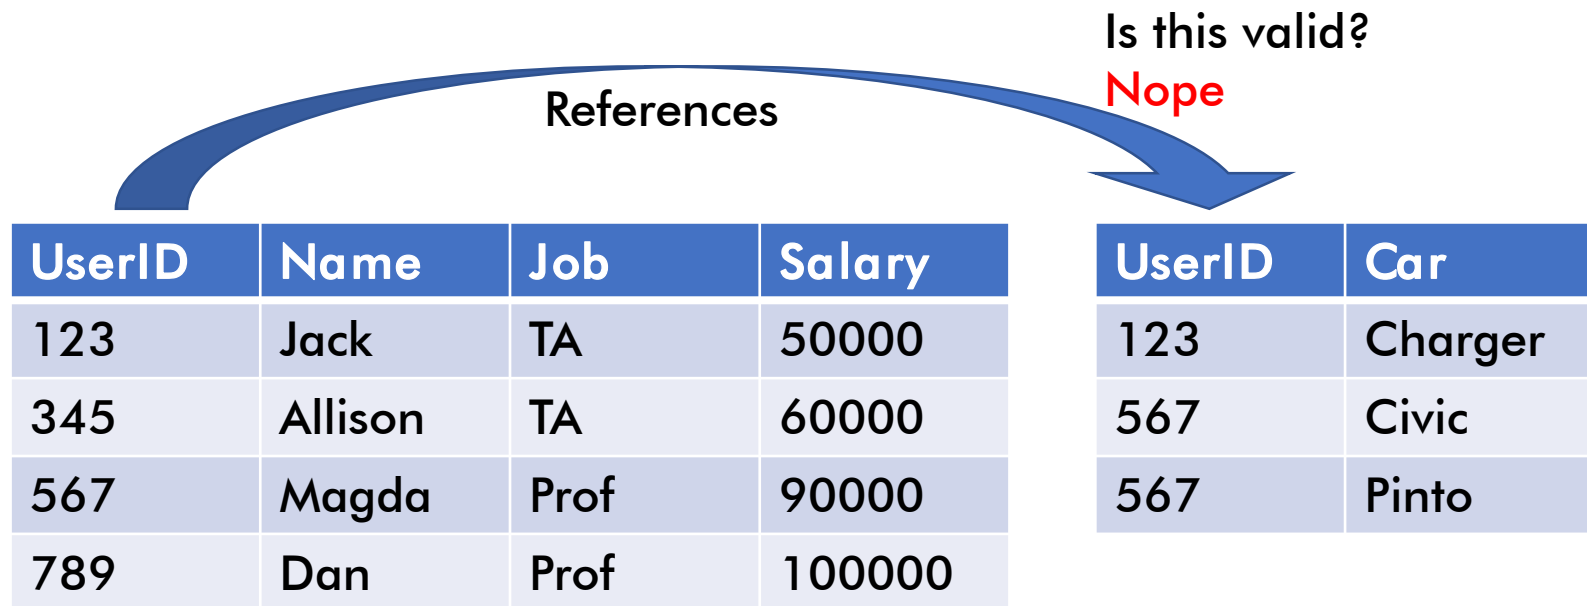
UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Foreign Keys

Foreign Key

A **Foreign Key** is one or more attributes that uniquely identify a row in *another table*.



Foreign Keys

```
CREATE TABLE Payroll (  
  UserID INT PRIMARY KEY,  
  Name TEXT,  
  Job TEXT,  
  Salary INT);
```

Payroll(UserId, Name, Job, Salary)

```
CREATE TABLE Regist (  
  UserID INT,  
  Car TEXT);
```

Regist(UserId, Car)

Foreign Keys

```
CREATE TABLE Payroll (  
    UserID INT PRIMARY KEY,  
    Name TEXT,  
    Job TEXT,  
    Salary INT);
```

```
CREATE TABLE Regist (  
    UserID INT REFERENCES Payroll,  
    Car TEXT);
```

Payroll(UserId, Name, Job, Salary)

Regist(UserId, Car)

The Relational Model Revisited

- **More complete overview of the Relational Model:**
 - Database → collection of tables
 - All tables are flat
 - Keys uniquely ID rows
 - Foreign keys act as a “semantic pointer”
 - **Physical data independence**

- Foreign keys *describe* a relationship between tables
- Joins *realize* combinations of data across relationships

Inner Joins

- Bread and butter of SQL queries
 - “Inner join” is often interchangeable with just “join”

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

How do we
algorithmically
get our results?

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics


Payroll



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
------	-----

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```


Nested-Loop Semantics

Payroll



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist




UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics


Payroll



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics


Payroll



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```


Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```


Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto


Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics


Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



Regist

UserID	Car
123	Charger
567	Civic
567	Pinto



Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Nested-Loop Semantics

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Magda	Civic
Magda	Pinto

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

Inner Joins

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

Explicit

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

Implicit

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

Inner Joins: RA

Payroll


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

2-arg Cartesian product op (Symbol "x")
Take all pairs of tuples



Cartesian product

UserID	Name	Job	Salary	UserID	Car
123	Jack	TA	50000	123	Charger
123	Jack	TA	50000	567	Civic
123	Jack	TA	50000	567	Pinto
345	Allison	TA	60000	123	Charger
345	Allison	TA	60000	567	Civic
...

(4 x 3 = 12 rows total)

×

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

Inner Joins: RA

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

2-arg Join op (Symbol "bowtie")
Take Cartesian product & filter

$$\bowtie_{condition} = \begin{array}{c} \sigma_{condition} \\ | \\ \times \end{array}$$

$$\begin{array}{c} \Pi_{P.Name, R.Car} \\ | \\ \bowtie_{P.UserID=R.UserID} \\ \swarrow \quad \searrow \\ Payroll\ P \quad Regist\ R \end{array}$$

Outer Joins

Now I want to include everyone, even if they don't drive.

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

Outer Joins

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

NULL is a value placeholder. Depending on context, it may mean unknown, not applicable, etc.

Name	Car
Jack	Charger
Allison	NULL
Magda	Civic
Magda	Pinto
Dan	NULL

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

Outer Joins: RA

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	40000
567	Magda	TA	40000
789	Dan	TA	40000

2-arg Left Outer Join op
(Symbol "bowtie with left edge")
Take Join + pair non-matching
entries on left with NULLs on right

UserID	Car
123	Charger
567	Civic
567	Pinto

Name	Car
Jack	Charger
Allison	NULL
Magda	Civic
Magda	Pinto
Dan	NULL

$\Pi_{P.Name, R.Car}$
 $\bowtie_{P.UserID = R.UserID}$
 $Payroll\ P \quad Regist\ R$

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

Outer Joins

- **LEFT OUTER JOIN**



- All rows in left table are preserved

- **RIGHT OUTER JOIN**



- All rows in right table are preserved

- **FULL OUTER JOIN**



- All rows are preserved

Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic';
```

Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto';
```

Will this work?

Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R.Car
  FROM Payroll AS P, Regist AS R
 WHERE P.UserID = R.UserID AND
        R.Car = 'Civic' AND
        R.Car = 'Pinto';
```

Will this work?
**Nope, empty set
is returned**

Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```


Self Joins

Find all people who drive a Civic and Pinto

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

All pairs of cars a person can drive

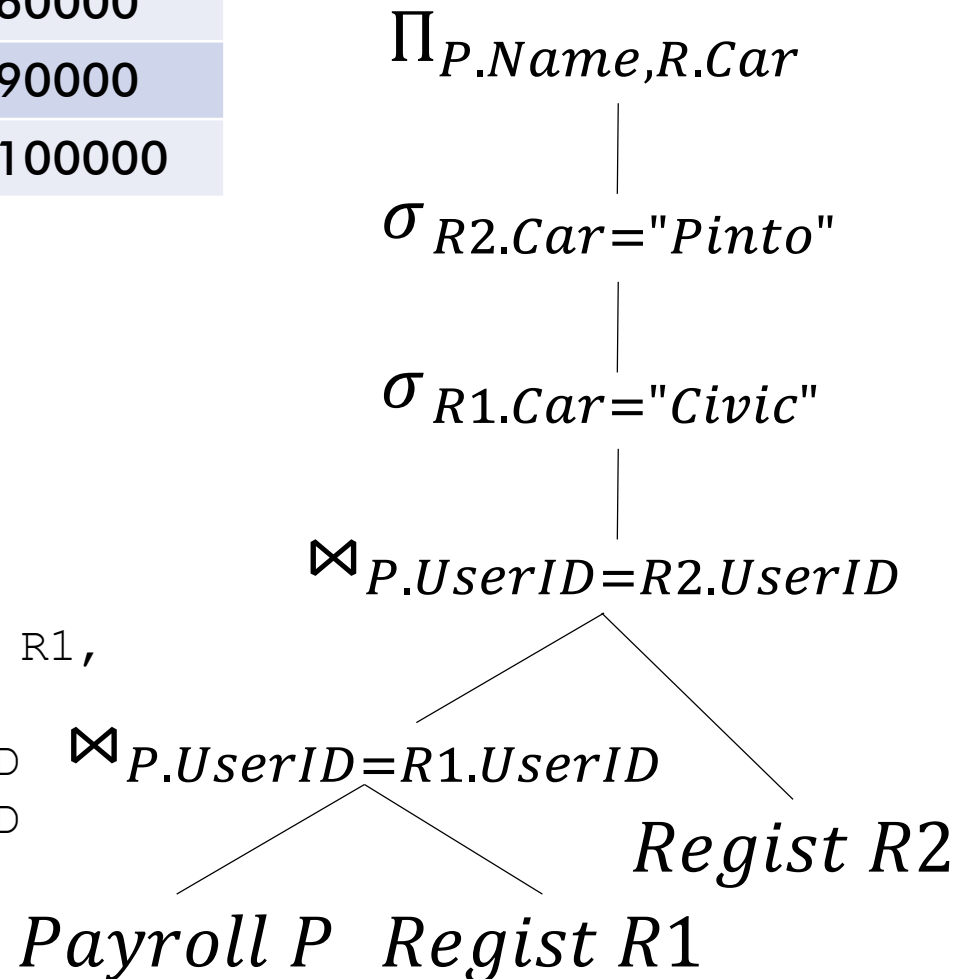
```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

Self Joins

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1,
      Regist AS R2
WHERE P.UserID = R1.UserID AND
        P.UserID = R2.UserID AND
        R1.Car = 'Civic' AND
        R2.Car = 'Pinto';
```

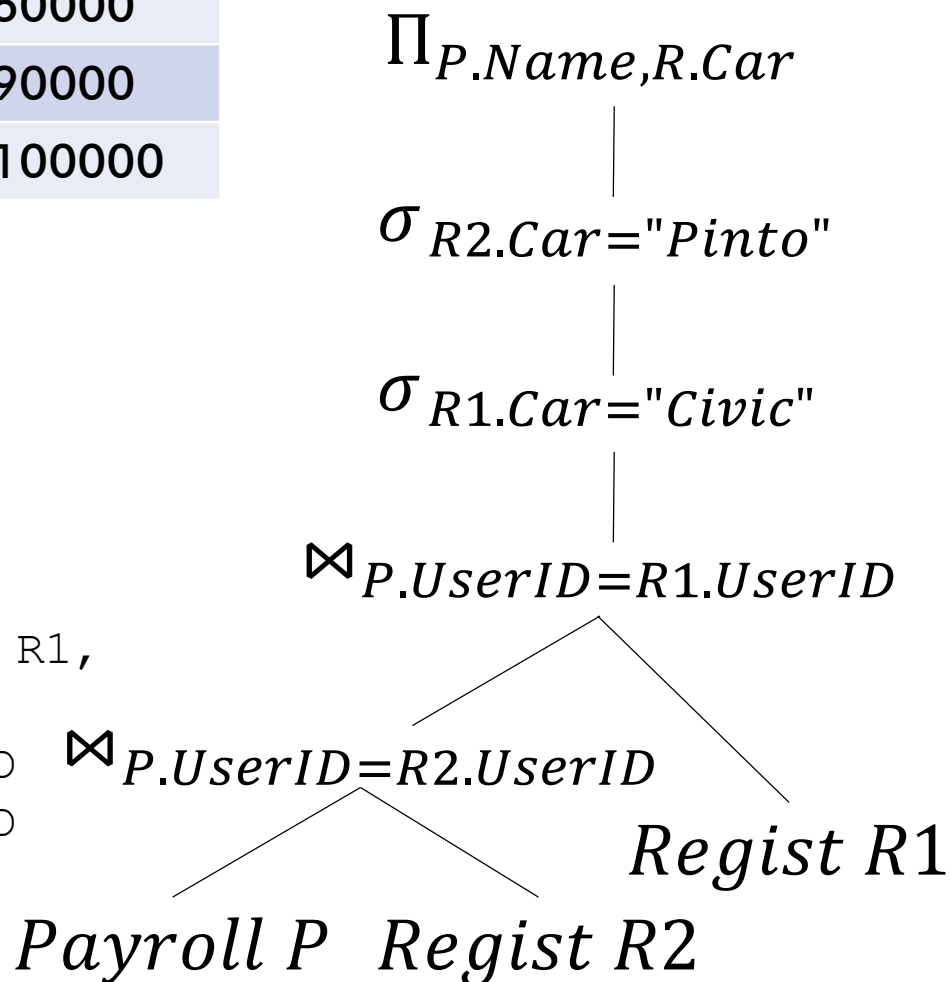


Self Joins (Equivalent RA)

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1,
      Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

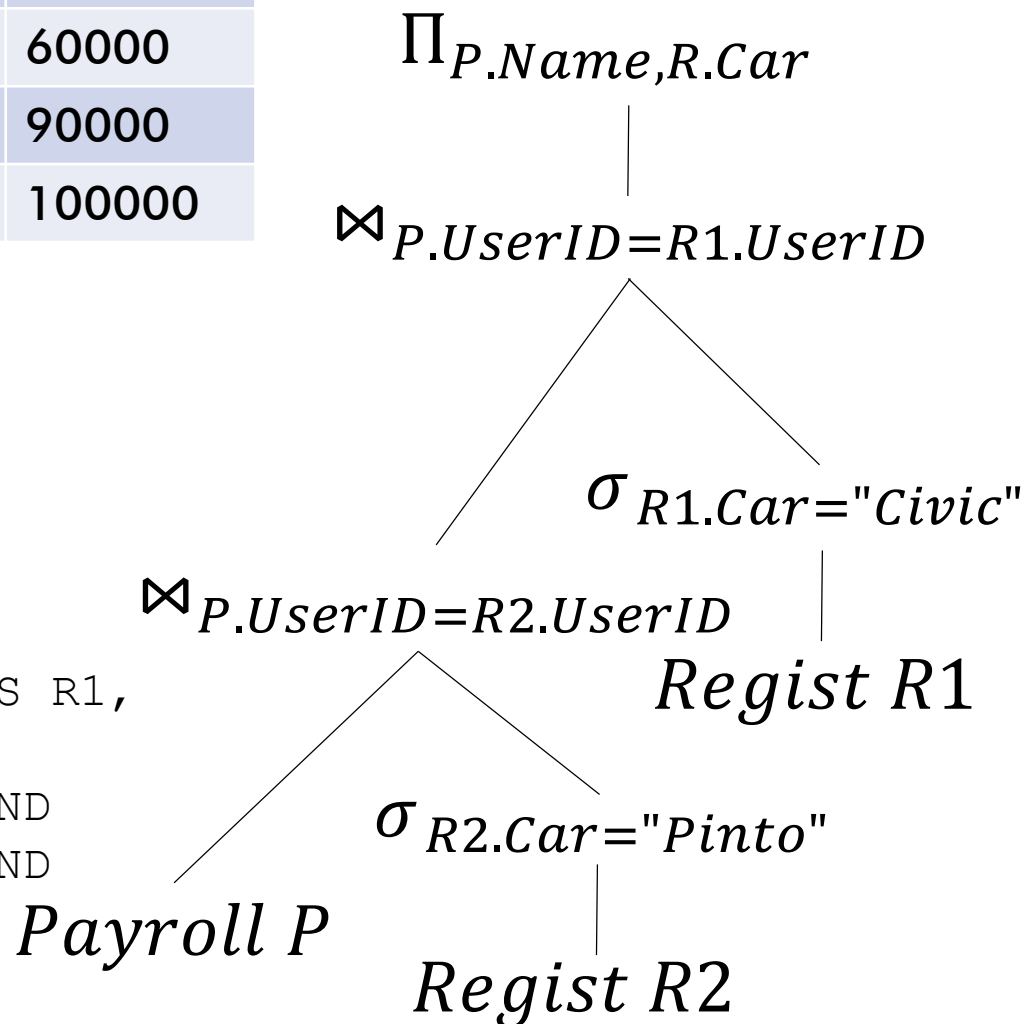


Self Joins (Equivalent RA)

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1,
      Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```



Takeaways

- We can describe relationships between tables with keys and foreign keys
- Different joining techniques can be used to achieve particular goals
- RA plans can be rearranged equivalently
- Our SQL toolbox is growing!
 - Not just reading and filtering data anymore
 - Starting to answer complex questions