

# Introduction to Data Management

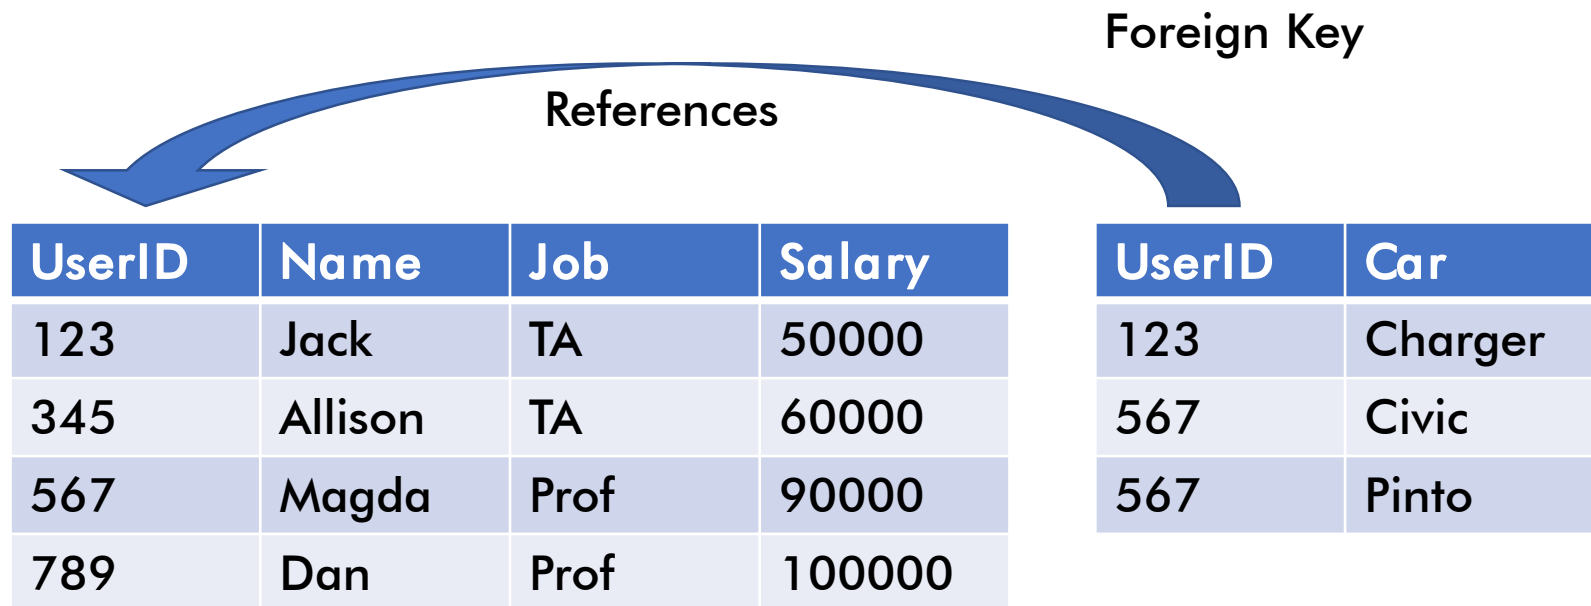
*Feeling aggregated?*

Shana Hutchison

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle


# Recap – Keys and Foreign Keys

- Modeling multiple tables in the same database
  - Keys and foreign keys



# Recap – Joins

- Join combine data across tables
  - Nested-loop semantics
  - Filtered Cartesian product semantics
  - Inner join (most common)
  - Outer joins preserve non-matching tuples
  - Self join pattern



# Actionable Results

- **Summaries of data help make decisions and succinctly convey information**
  - “How popular is this anime?” → COUNT
  - “Do I spend too much on tea?” → SUM
  - “Am I being ripped off by this dealer?” → AVG
  - “Who raised the most money for charity?” → MAX
  - “What is the cheapest Greek yogurt?” → MIN

# Actionable Results

- **Summaries of data help make decisions and succinctly convey information**

- **SELECT COUNT(\*) FROM AnimeVideoViews ...**
- **SELECT SUM(cost) FROM TeaReceipts ...**
- **SELECT AVG(price) FROM CarDealers ...**
- **SELECT MAX(score) FROM Donations ...**
- **SELECT MIN(price) FROM YogurtStores ...**




**AGG(attr) → computes AGG over non-NULL values**  
**AGG(DISTINCT attr) is also possible**

**Watch out for NULLs!**

# Actionable Results

- **Summaries of data help make decisions and succinctly convey information**

- **SELECT COUNT(\*) FROM AnimeVideoViews ...**
- SELECT SUM(cost) FROM TeaReceipts ...
- SELECT AVG(price) FROM CarDealers ...
- SELECT MAX(score) FROM Donations ...
- SELECT MIN(price) FROM YogurtStores ...



**COUNT(\*) → # of rows  
regardless of NULL**

**Watch out for NULLs!**

# Aggregation Semantics

What am I aggregating over in a **SELECT-FROM-WHERE** query?

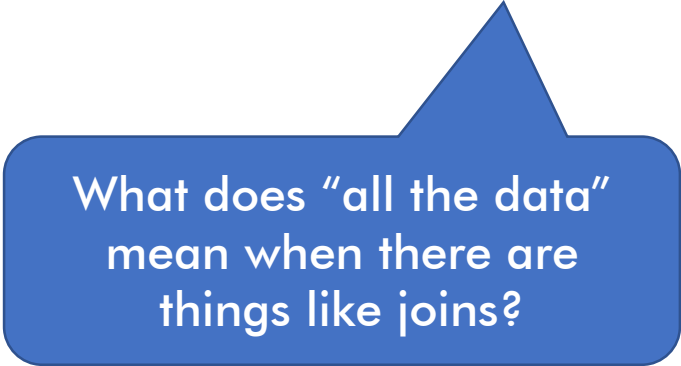
Intuitively: “all the data”



# Aggregation Semantics

What am I aggregating over in a SELECT-FROM-WHERE query?

Intuitively: “all the data”



What does “all the data”  
mean when there are  
things like joins?



# Aggregation Semantics

What am I aggregating over in a **SELECT-FROM-WHERE** query?

```
SELECT  AVG(P.Salary)
FROM    Payroll AS P, Regist AS R
WHERE    P.UserID = R.UserID;
```

**Payroll**


UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Regist**

UserID	Car
123	Charger
567	Civic
567	Pinto

# Aggregation Semantics

```
SELECT AVG(P.Salary)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

  $P.UserID = R.UserID$

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Aggregation Semantics

```
SELECT AVG(P.Salary)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

P.UserID	P.Name	P.Job	P.Salary	R.UserID	R.Car
123	Jack	TA	50000	123	Charger
567	Magda	Prof	90000	567	Civic
567	Magda	Prof	90000	567	Pinto

$\bowtie P.UserID=R.UserID$

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Aggregation Semantics

```
SELECT AVG(P.Salary)
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

$\gamma_{AVG(P.Salary)}$

P.UserID	P.Name	P.Job	P.Salary	R.UserID	R.Car
123	Jack	TA	50000	123	Charger
567	Magda	Prof	90000	567	Civic
567	Magda	Prof	90000	567	Pinto

$\bowtie P.UserID=R.UserID$

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Aggregation Semantics

```
SELECT AVG(P.Salary)
  FROM Payroll AS P, Regist AS R
 WHERE P.UserID = R.UserID;
```

**AVG(P.Salary)**

76666

$\gamma_{AVG(P.Salary)}$

P.UserID	P.Name	P.Job	P.Salary	R.UserID	R.Car
123	Jack	TA	50000	123	Charger
567	Magda	Prof	90000	567	Civic
567	Magda	Prof	90000	567	Pinto

$\bowtie_{P.UserID=R.UserID}$

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto

# Aggregation Semantics

```
SELECT AVG(P.Salary)
  FROM Payroll AS P, Regist AS R
 WHERE P.UserID = R.UserID;
```

1-arg Aggregation op (Greek "gamma")  
Compute aggregates,  
grouping by non-aggregates

$\gamma_{AVG(P.Salary)}$

Like projection,  
only keeps  
listed attributes

$\bowtie P.UserID=R.UserID$

*Payroll P*

*Regist R*

- SQL allows you to specify what groups your query operates over
  - Sometimes a “whole-table” aggregation is too coarse-grained
  - We can partition our data based on **matching attribute values**

# Grouping

- SQL allows you to specify what groups your query operates over
  - Sometimes a “whole-table” aggregation is too coarse-grained
  - We can partition our data based on **matching attribute values**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

...

**GROUP BY** Job

...



# Grouping

- SQL allows you to specify what groups your query operates over
  - Sometimes a “whole-table” aggregation is too coarse-grained
  - We can partition our data based on **matching attribute values**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

...

**GROUP BY** Job

...

# Grouping Example

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Grouping Example

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)
TA	60000
Prof	100000

# Grouping on Multiple Attributes

```
SELECT Name, MAX (Salary)
FROM Payroll
GROUP BY Job, Name
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Name	Salary
Jack	50000
Allison	60000
Magda	90000
Dan	100000

# Grouping on Multiple Attributes

```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job, Name
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
424	Jack	TA	55000

Name	Salary
Jack	55000
Allison	60000
Magda	90000
Dan	100000

# Filtering Groups with HAVING

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Filtering Groups with HAVING

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)
Prof	100000

How is aggregation processed internally?

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



# Aggregation RA

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

Example RA syntax  
for creating aliases  
maxSal, minSal

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

Job	maxSal	minSal
Prof	100000	90000

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$

Job	maxSal	minSal
Prof	100000	90000

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
  FROM Payroll
 GROUP BY Job
HAVING MIN(Salary) > 80000
```

Job	maxSal
Prof	100000

$\Pi_{Job, maxSal}$

Job	maxSal	minSal
Prof	100000	90000

$\sigma_{minSal > 80000}$

Job	maxSal	minSal
TA	60000	50000
Prof	100000	90000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

UserID	Name	Job	Salary
...	...	...	...

# Aggregation RA

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$

HAVING is a  
selection operator  
after aggregation


$\sigma_{minSal > 80000}$

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

*Payroll P*




# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
-----	-------------	-------------


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```




# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	50000	50000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```




# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	50000	50000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
  insert into a dictionary D
    row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
  if (row.MIN(Salary) > 80000)
    output (row.Job, row.MAX(Salary))
```




# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```




# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```



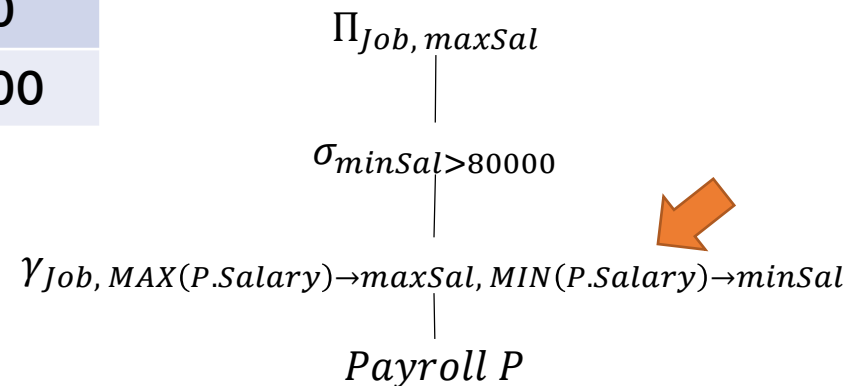
# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```


Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	90000	90000



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```




# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	90000	90000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
  insert into a dictionary D
    row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
  if (row.MIN(Salary) > 80000)
    output (row.Job, row.MAX(Salary))
```



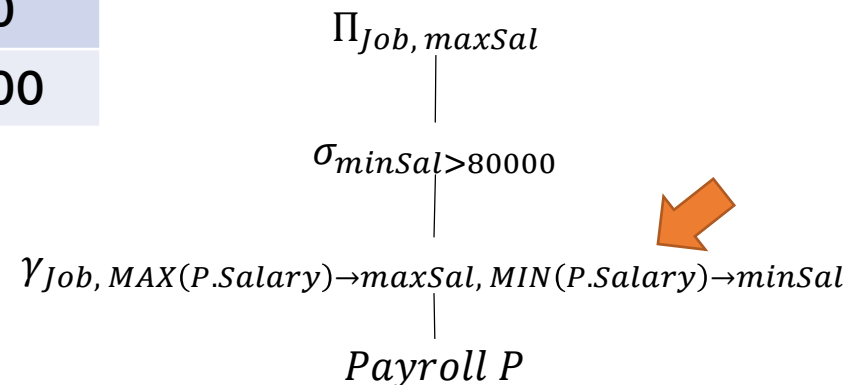
# Agg: Nested-Loop Semantics




UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```





# Agg: Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$

$\sigma_{minSal > 80000}$


$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

Payroll P



Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000

```
for each row in Payroll:
  insert into a dictionary D
    row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
  if (row.MIN(Salary) > 80000)
    output (row.Job, row.MAX(Salary))
```



# Agg: Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$  ←

$\sigma_{minSal > 80000}$  ←

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$


Payroll P



Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000

Job	MAX(Salary)
-----	-------------

```
for each row in Payroll:
  insert into a dictionary D
    row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
  if (row.MIN(Salary) > 80000)
    output (row.Job, row.MAX(Salary))
```



# Agg: Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$  ←

$\sigma_{minSal > 80000}$  ←


$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

Payroll P

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000

Job	MAX(Salary)
-----	-------------

```
for each row in Payroll:
  insert into a dictionary D
    row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
  if (row.MIN(Salary) > 80000)
    output (row.Job, row.MAX(Salary))
```



# Agg: Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$  ←

$\sigma_{minSal > 80000}$  ←

$\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$

Payroll P

Job	MAX(Salary)
Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000

for each row in Payroll:


    insert into a dictionary D

        row.Job => MAX(row.Salary), MIN(row.Salary)

for each row in D:

    if (row.MIN(Salary) > 80000)

        output (row.Job, row.MAX(Salary))

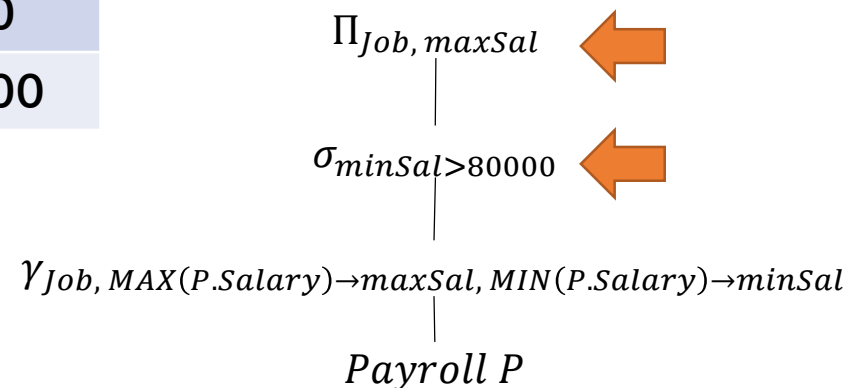


# Agg: Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```



Job	MAX(Salary)
Prof	100000



```
for each row in Payroll:
    insert into a dictionary D
        row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
    if (row.MIN(Salary) > 80000)
        output (row.Job, row.MAX(Salary))
```



# Agg: Nested-Loop Semantics

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Job	MAX(Salary)	MIN(Salary)
TA	60000	50000
Prof	100000	90000


```
SELECT Job, MAX(Salary)
FROM Payroll
GROUP BY Job
HAVING MIN(Salary) > 80000
```

$\Pi_{Job, maxSal}$   
|  
 $\sigma_{minSal > 80000}$   
|  
 $\gamma_{Job, MAX(P.Salary) \rightarrow maxSal, MIN(P.Salary) \rightarrow minSal}$   
|  
Payroll P

Job	MAX(Salary)
Prof	100000

```
for each row in Payroll:
  insert into a dictionary D
    row.Job => MAX(row.Salary), MIN(row.Salary)
for each row in D:
  if (row.MIN(Salary) > 80000)
    output (row.Job, row.MAX(Salary))
```

# SQL and RA Vocab Summary



# SQL and RA Vocab Summary

**SELECT** (**DISTINCT**) ...  
**FROM** ...  
**WHERE** ...  
**GROUP BY** ...  
**HAVING** ...  
**ORDER BY** ...

$\delta$

$\Pi$

$\tau$

$\sigma$

$\gamma$

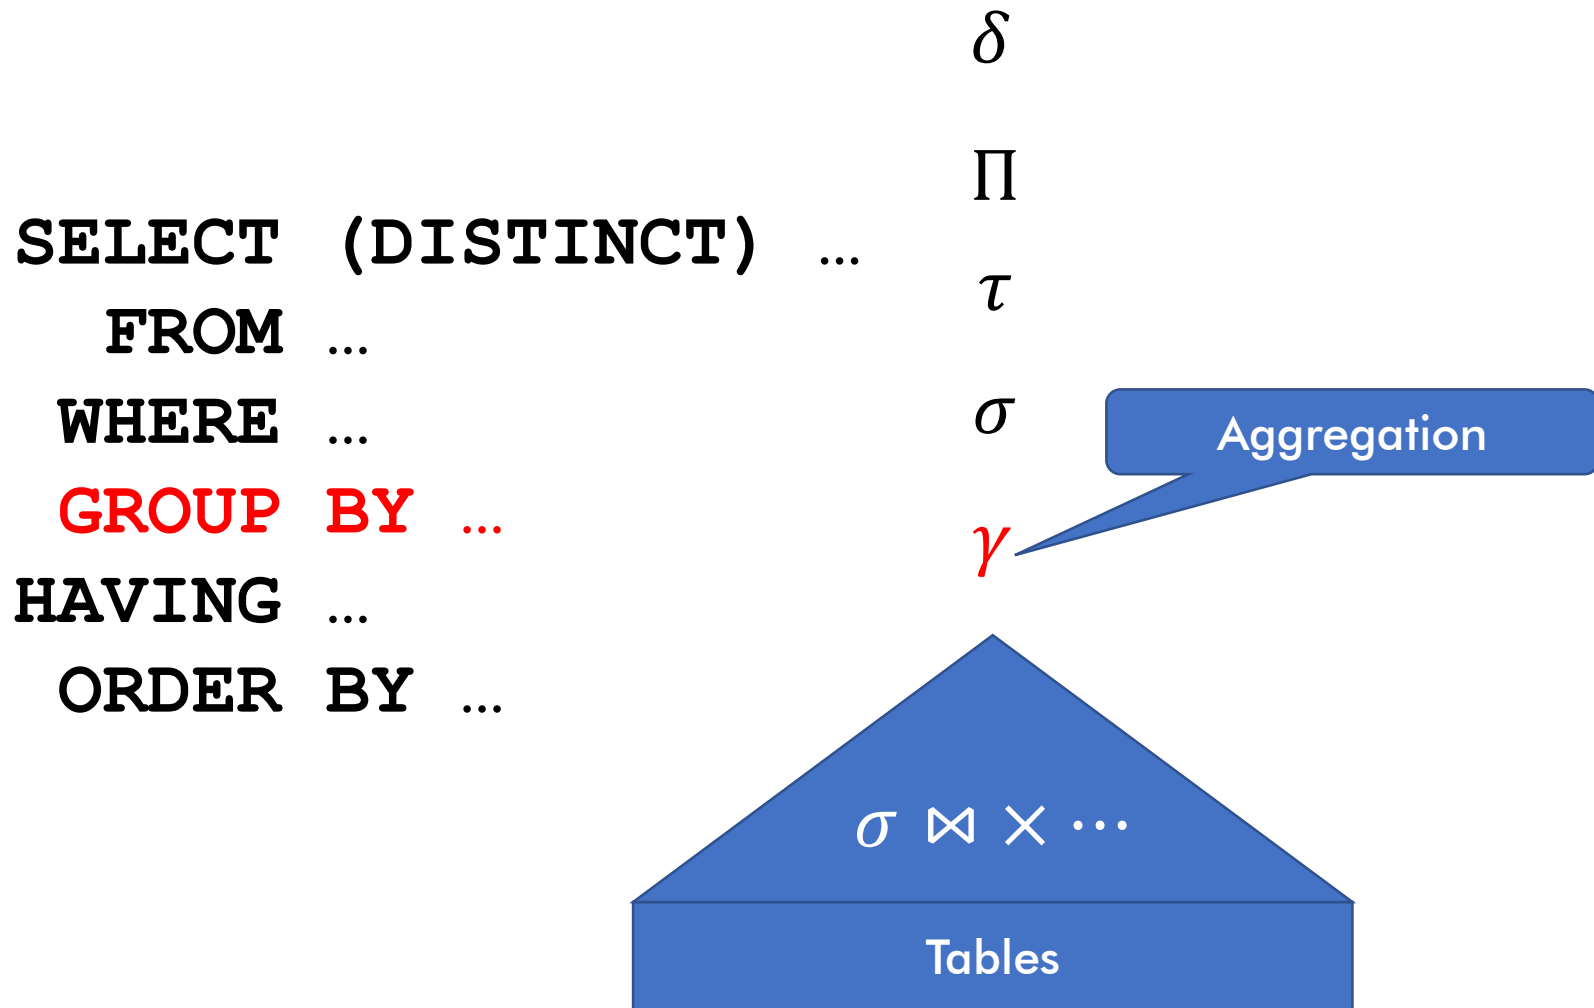
Selection  
Join  
Cartesian Product

$\sigma \bowtie \times \dots$

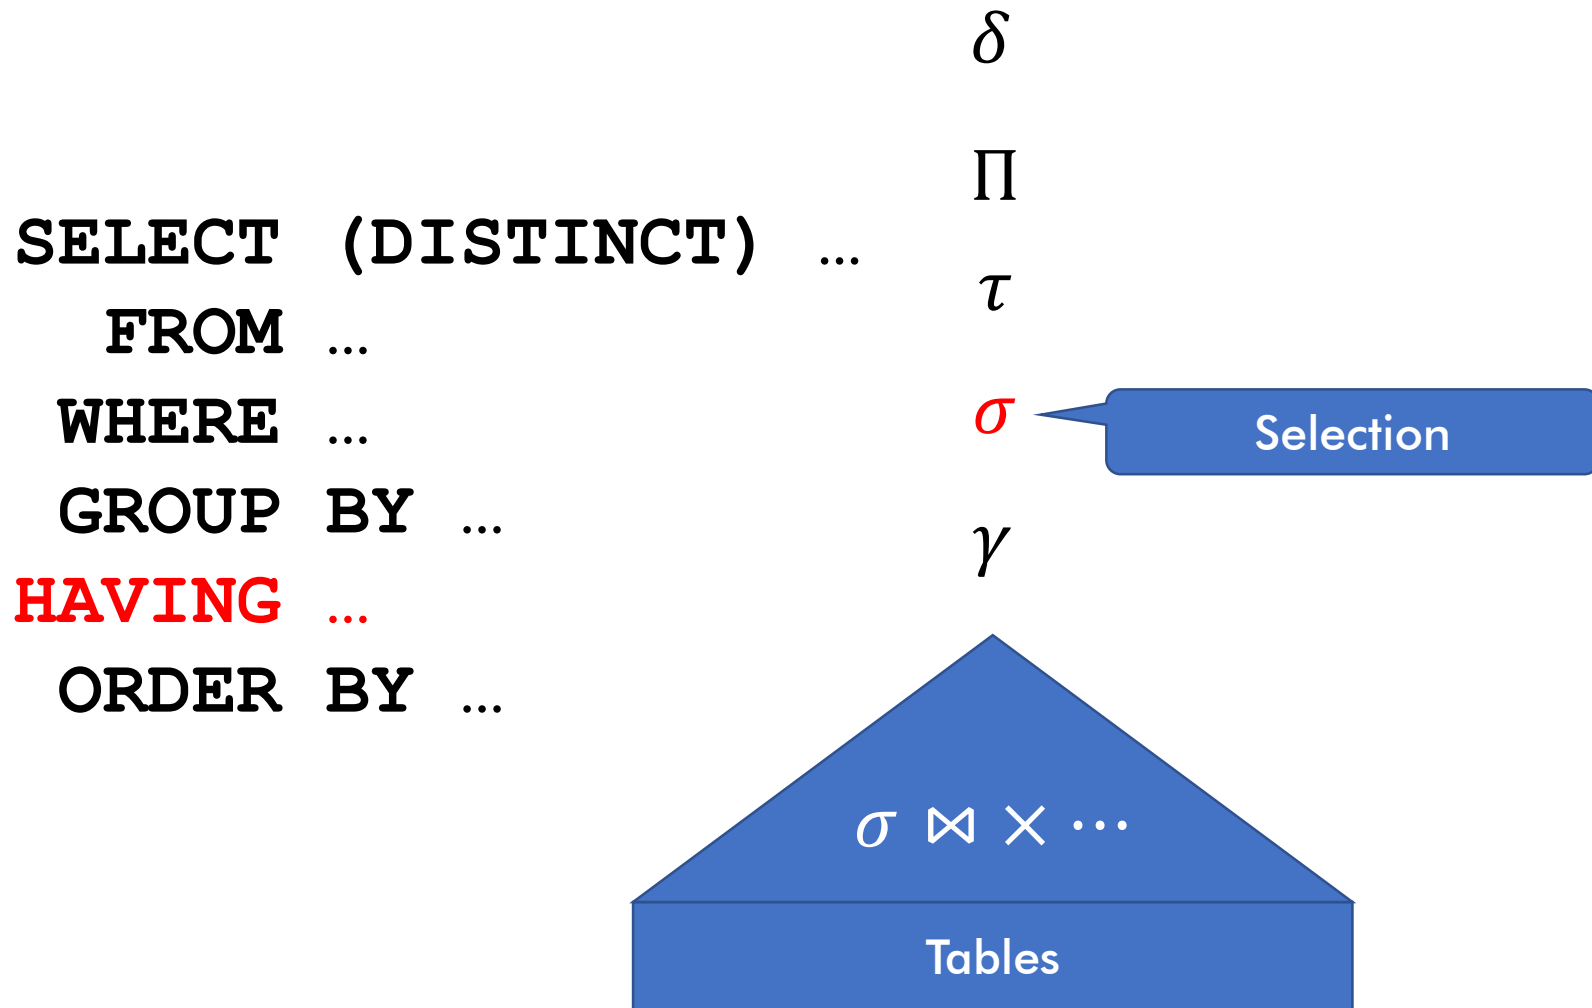
Tables



# SQL and RA Vocab Summary



# SQL and RA Vocab Summary



# SQL and RA Vocab Summary

SELECT (DISTINCT) ...  
FROM ...  
WHERE ...  
GROUP BY ...  
HAVING ...  
ORDER BY ...

$\delta$

$\Pi$

$\tau$

$\sigma$

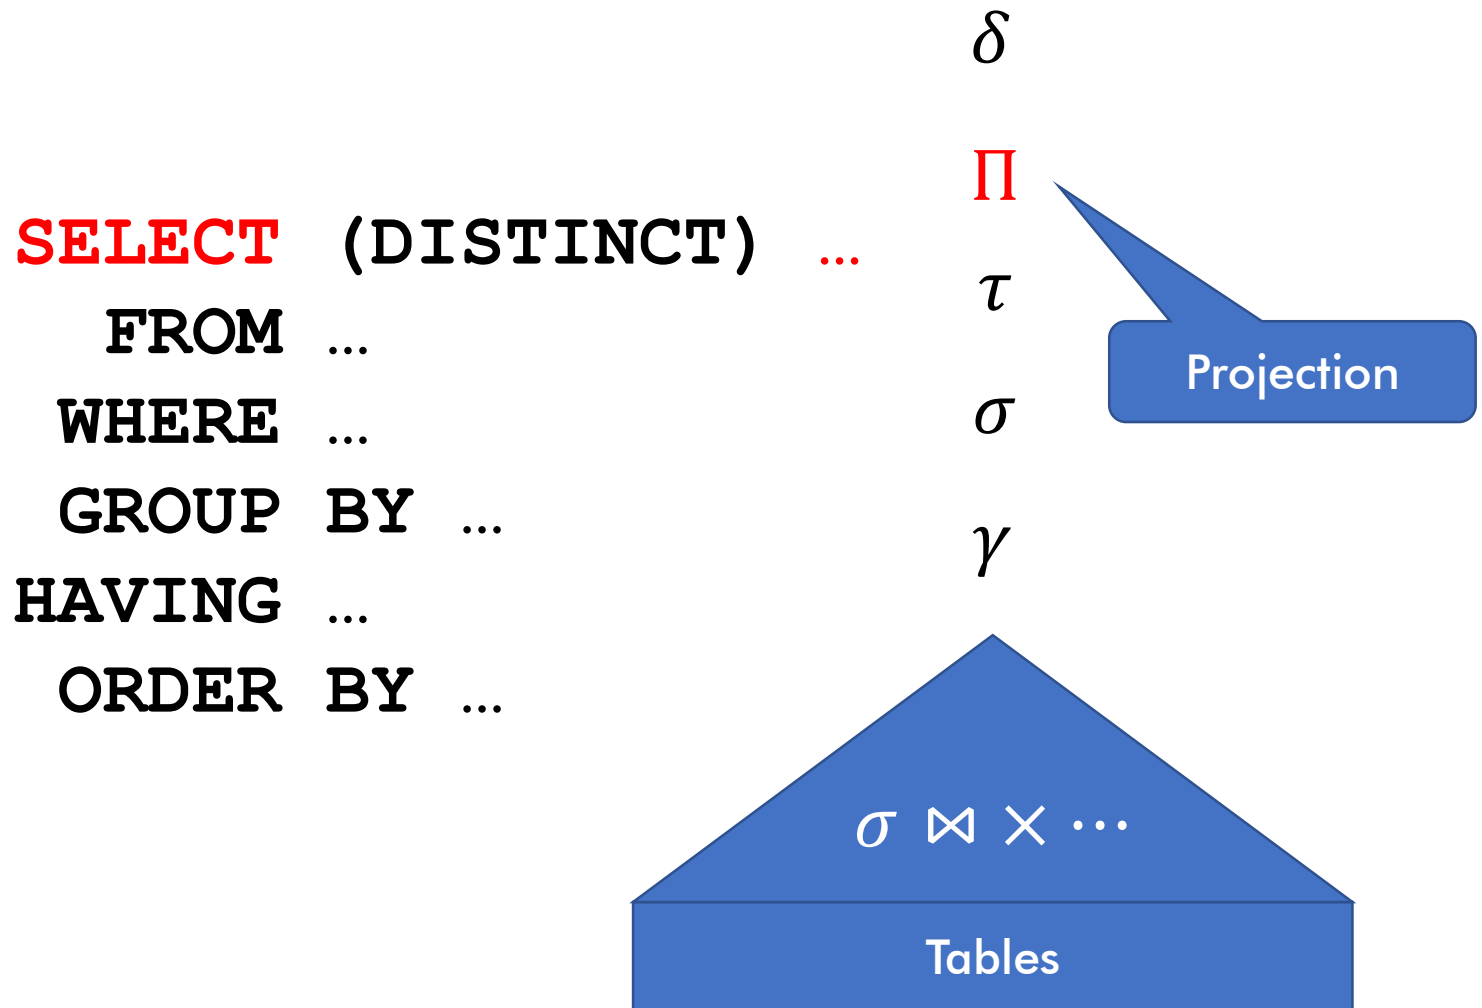
$\gamma$

Sorting

$\sigma \bowtie \times \dots$

Tables

# SQL and RA Vocab Summary



# SQL and RA Vocab Summary

**SELECT**    **(DISTINCT)**    ...  
      **FROM**    ...  
      **WHERE**    ...  
      **GROUP BY**    ...  
**HAVING**    ...  
      **ORDER BY**    ...

 $\delta$  $\Pi$  $\tau$  $\sigma$  $\gamma$ 

Duplicate Elim

 $\sigma \bowtie \times \dots$ 

Tables

## FWGHOS™

**SELECT (DISTINCT) ...**  
**FROM ...**  
**WHERE ...**  
**GROUP BY ...**  
**HAVING ...**  
**ORDER BY ...**

$\delta$

$\Pi$

$\tau$

$\sigma$

$\gamma$

$\sigma \bowtie \times \dots$

Tables

# The Witnessing Problem

- Also known as **argmax/argmin**
- Ex: Return the person with the highest salary for each job type

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# The Witnessing Problem

- Also known as **argmax/argmin**
- Ex: Return the person with the highest salary for each job type

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job
```

Easy, right?



# The Witnessing Problem

- Also known as **argmax/argmin**
- Ex: Return the person with the highest salary for each job type

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
242	Gibbs	TA	60000



```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job
```

# The Witnessing Problem

- Also known as argmax

- Ex: Return the highest salary for each job

Failed to execute query. Error: Column 'Payroll.name' is invalid in the select list because it is not contained in either an aggregate function or the GROUP BY clause.

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
242	Gibbs	TA	60000



```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job
```

SELECT, HAVING, ORDER BY  
can only use GROUP BY  
attributes or aggregates

# The Witnessing Problem

- Also known as `argmax()`

- Ex: Return the highest salary for each job

Failed to execute query. Error: Column 'Payroll.name' is invalid in the select list because it is not contained in either an aggregate function or the GROUP BY clause.

UserID	Name	Job	Salary
123	Jack	TA	80000
345	Allison	TA	80000
567	Magda	Prof	90000
789	Dan	Prof	100000
242	Gibbs	TA	60000

This is a problem even for 'Prof' where it seems like Dan is the clear answer. SQL Standard says it *could be ambiguous* in general, so it is never allowed

SQLite allows it... returns a random name



```
SELECT Name, MAX(Salary)
FROM Payroll
GROUP BY Job
```

SELECT, HAVING, ORDER BY  
can only use GROUP BY  
attributes or aggregates

# The Witnessing Problem

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000
242	Gibbs	TA	60000

Return the person with the highest salary for each job type

How do we witness the maxima for a group?

**Discuss!**

Conceptual ideas are great

# The Witnessing Problem

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000
242	Gibbs	TA	60000	60000

Return the person with the highest salary for each job type

Main idea: we need to join the respective maxima to each row

# The Witnessing Problem

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000
242	Gibbs	TA	60000	60000

Return the person with the highest salary for each job type

Main idea: we need to join the respective maxima to each row

# The Witnessing Problem

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000
242	Gibbs	TA	60000	60000

Return the person with the highest salary for each job type

```
SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```

# The Witnessing Problem

Result of the Join

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
123	Jack	TA	50000	242	Gibbs	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	242	Gibbs	TA	60000
242	Gibbs	TA	60000	123	Jack	TA	50000
242	Gibbs	TA	60000	345	Allison	TA	60000
242	Gibbs	TA	60000	242	Gibbs	TA	60000
567	Magda	Prof	70000	567	Magda	Prof	70000
567	Magda	Prof	70000	789	Dan	Prof	80000
789	Dan	Prof	80000	789	Dan	Prof	80000
789	Dan	Prof	80000				

```

SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
    
```



# The Witnessing Problem

## Grouping

P1

P2

UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
123	Jack	TA	50000	242	Gibbs	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	242	Gibbs	TA	60000
242	Gibbs	TA	60000	123	Jack	TA	50000
242	Gibbs	TA	60000	345	Allison	TA	60000
242	Gibbs	TA	60000	242	Gibbs	TA	60000
567	Magda	Prof		<b>SELECT</b> P1.Name, MAX(P2.Salary) <b>FROM</b> Payroll AS P1, Payroll AS P2 <b>WHERE</b> P1.Job = P2.Job <b>GROUP BY</b> P2.Job, P1.Salary, P1.Name <b>HAVING</b> P1.Salary = MAX(P2.Salary)			
567	Magda	Prof					
789	Dan	Prof					
789	Dan	Prof					

# The Witnessing Problem

Max within groups							
P1				P2			
UserID	Name	Job	Salary	UserID	Name	Job	Salary
123	Jack	TA	50000	123	Jack	TA	50000
123	Jack	TA	50000	345	Allison	TA	60000
123	Jack	TA	50000	242	Gibbs	TA	60000
345	Allison	TA	60000	123	Jack	TA	50000
345	Allison	TA	60000	345	Allison	TA	60000
345	Allison	TA	60000	242	Gibbs	TA	60000
242	Gibbs	TA	60000	123	Jack	TA	50000
242	Gibbs	TA	60000	345	Allison	TA	60000
242	Gibbs	TA	60000	242	Gibbs	TA	60000
567	Magda						
567	Magda						
789	Dan						
789	Dan						

**SELECT** P1.Name, MAX(P2.Salary)

**FROM** Payroll AS P1, Payroll AS P2

**WHERE** P1.Job = P2.Job

**GROUP BY** P2.Job, P1.Salary, P1.Name

**HAVING** P1.Salary = MAX(P2.Salary)

# The Witnessing Problem

Cut down to groups

Name	Salary	Job	MAX(Salary)
Jack	50000	TA	60000
Allison	60000	TA	60000
Gibbs	60000	TA	60000
Magda	90000	Prof	100000
Dan	100000	Prof	100000

```
SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
 GROUP BY P2.Job, P1.Salary, P1.Name
 HAVING P1.Salary = MAX(P2.Salary)
```

# The Witnessing Problem

HAVING filter

Name	Salary	Job	MAX(Salary)
Jack	50000	TA	60000
Allison	60000	TA	60000
Gibbs	60000	TA	60000
Magda	90000	Prof	100000
Dan	100000	Prof	100000



```
SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
 GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```



# The Witnessing Problem

SELECT projection



Name	Salary	Job	MAX(Salary)
Allison	60000	TA	60000
Gibbs	60000	TA	60000
Dan	100000	Prof	100000

```
SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
 GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```



# The Witnessing Problem

Output

Name	MAX(Salary)
Allison	60000
Gibbs	60000
Dan	100000

Warning: what if Allison and Gibbs had the same name?

→ Group by and select UserID to distinguish

```
SELECT P1.Name, MAX(P2.Salary)
  FROM Payroll AS P1, Payroll AS P2
 WHERE P1.Job = P2.Job
 GROUP BY P2.Job, P1.Salary, P1.Name
 HAVING P1.Salary = MAX(P2.Salary)
```

# Takeaways

- **FWGHOS™**
- **Combining techniques (aggregates and joins) allows you to answer complex questions (e.g. witnessing queries)**
- **Next week: simplifying with subqueries**