

# Introduction to Data Management

query in a query in a query in a ...

Shana Hutchison

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle

- Our SQL toolbox grows!
  - GROUP BY – grouping and aggregating
  - HAVING – post-group selection filter
- SQL “executed” in the following order
  1. FROM
  2. WHERE
  3. GROUP BY
  4. HAVING
  5. ORDER BY
  6. Select (DISTINCT)

# Recap – The Witnessing Problem

- A question pattern that asks for data associated with a maxima/minima of some value
  - Observed how to do it with grouping
  - “Self join” on values you find the maxima for
  - GROUP BY to deduplicate one side of the join
  - HAVING to compare values with respective maxima

- **Witnessing Problem**
  - (last lecture) with fancy self-join
  - building up subqueries from sub-plans
  - WITH clause to abstract useful intermediate result
- **Subquery mechanics**
  - Set/bag operations
  - FROM
  - SELECT
  - WHERE/HAVING
- **Decorrelation, equivalences along the way**
- **Universal quantification queries**

# The Witnessing Problem Simplified

- Wanted to join respective maxima
  - GROUP BY technique was interesting
  - Remember the suggestion last Friday that we **compute the maxima first then join & filter?**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Return the person (or people) with the highest salary for each job type

# The Witnessing Problem Simplified

- Wanted to join respective maxima
  - GROUP BY technique was interesting
  - Remember the suggestion last Friday that we **compute the maxima first then join & filter?**

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

Return the person (or people) with the highest salary for each job type

# The Witnessing Problem Simplified

```
SELECT Job,  
       MAX(Salary) AS maxima  
FROM Payroll  
GROUP BY Job
```

Job	maxima
TA	60000
Prof	100000



$\forall_{Job, MAX(P.Salary) \rightarrow maxima}$   
**Payroll**



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# The Witnessing Problem Simplified

```
SELECT Job,  
       MAX(Salary) AS maxima  
FROM Payroll  
GROUP BY Job
```

Job	maxima
TA	60000
Prof	100000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxima}$   
Payroll

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000



"Natural Join"  
Join on matching  
attributes

Okay, but how do we  
write SQL for this beast?

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



# The Witnessing Problem Simplified

②

```
SELECT Job,  
       MAX(Salary) AS maxima  
FROM Payroll  
GROUP BY Job
```

Job	maxima
TA	60000
Prof	100000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxima}$   
Payroll

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

```
SELECT *  
FROM ①, ②  
WHERE ①.Job = ②.Job
```

①

Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# The Witnessing Problem Simplified

```
SELECT Job,  
       MAX(Salary) AS maxima  
FROM Payroll  
GROUP BY Job
```

Job	maxima
TA	60000
Prof	100000

$\gamma_{Job, MAX(P.Salary) \rightarrow maxima}$   
Payroll

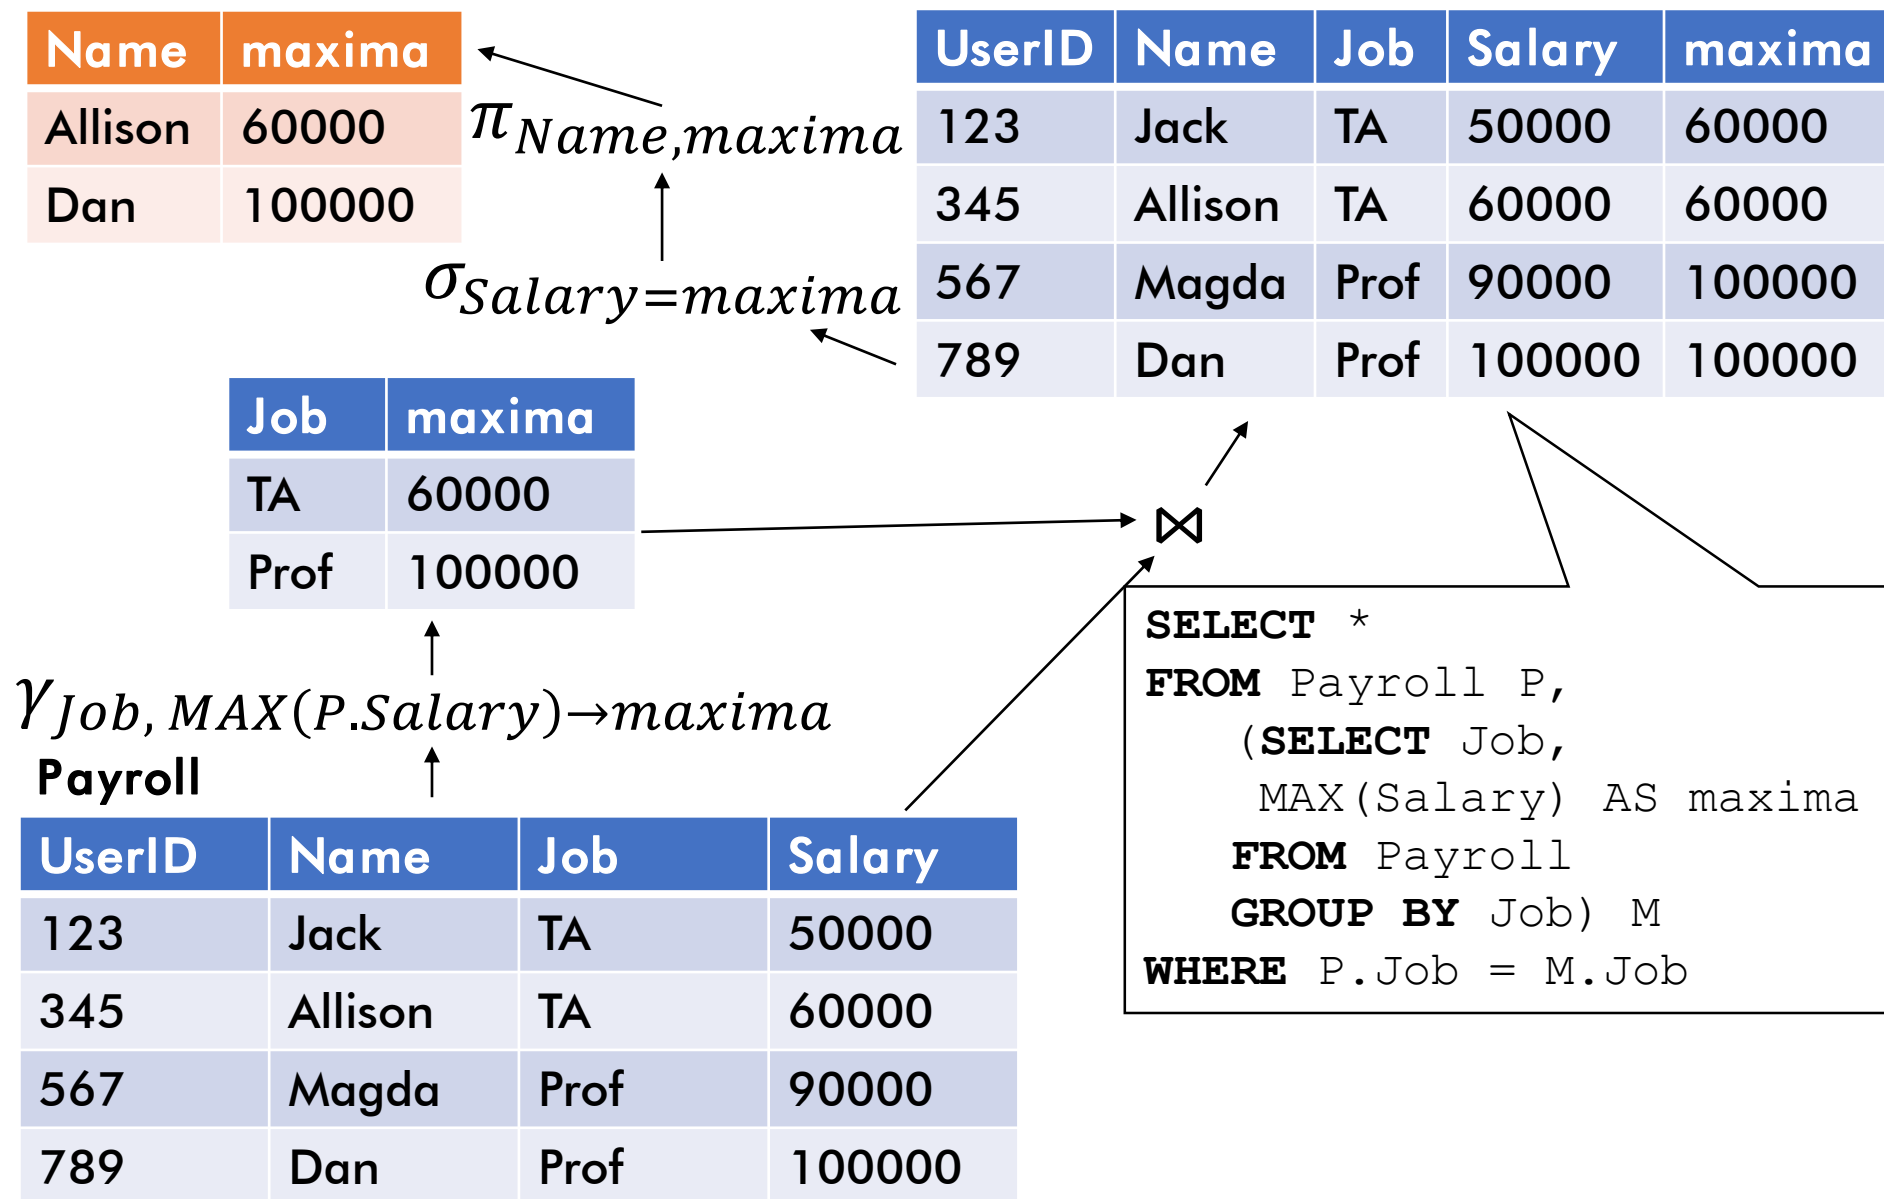
UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

```
SELECT *  
FROM Payroll P,  
     (SELECT Job,  
          MAX(Salary) AS maxima  
      FROM Payroll  
      GROUP BY Job) M  
WHERE P.Job = M.Job
```

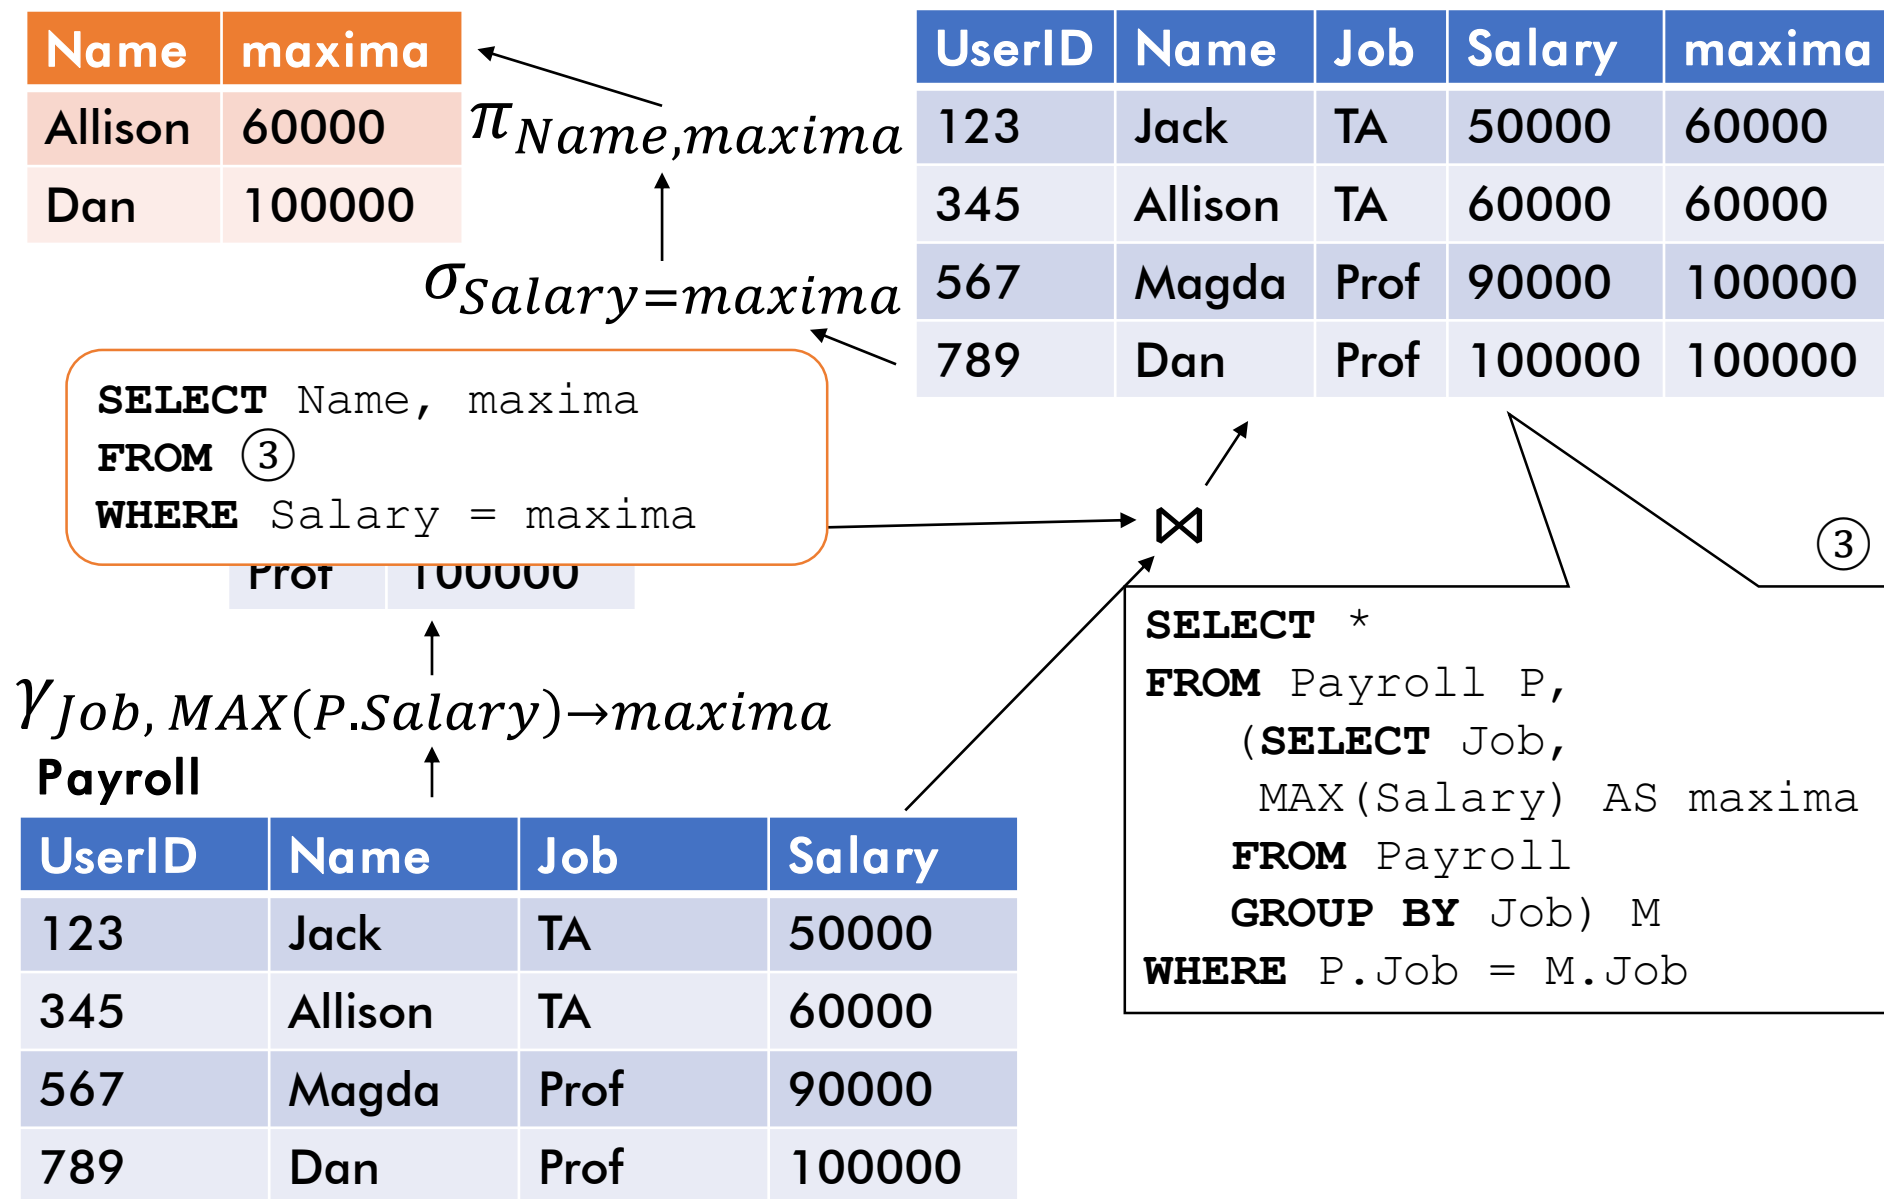
Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# The Witnessing Problem Simplified



# The Witnessing Problem Simplified



# The Witnessing Problem Simplified

Name	maxima
Allison	60000
Dan	100000

$\pi_{Name, maxima}$

$\sigma_{Salary=maxima}$

```

SELECT Name, maxima
FROM (SELECT *
      FROM Payroll P,
      (SELECT Job,
        MAX(Salary) AS maxima
      FROM Payroll
      GROUP BY Job) M
      WHERE P.Job = M.Job)
WHERE Salary = maxima
    
```

UserID	Name	Job	Salary	maxima
123	Jack	TA	50000	60000
345	Allison	TA	60000	60000
567	Magda	Prof	90000	100000
789	Dan	Prof	100000	100000

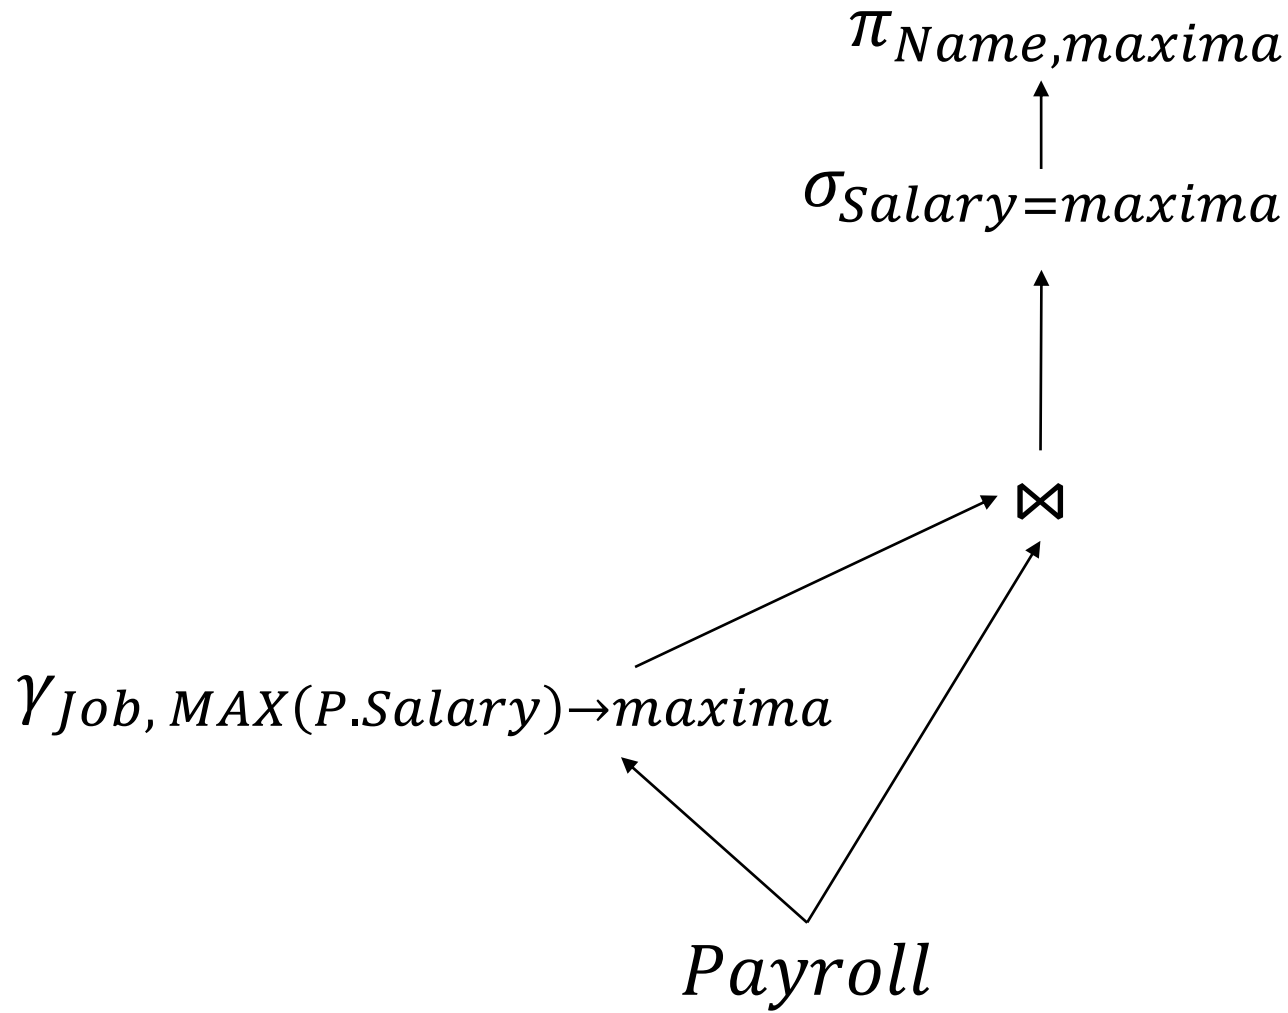
```

SELECT *
FROM Payroll P,
(SELECT Job,
  MAX(Salary) AS maxima
FROM Payroll
GROUP BY Job) M
WHERE P.Job = M.Job
    
```



123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# The Witnessing Problem Simplified





**Split your problem into sub-problems!**

**Solve them separately,  
then compose them!**

**Easier to debug a small query than a big one...**

# The Witnessing Problem [More] Simplified

- Hold up, too complicated!
- 1<sup>st</sup> & 2<sup>nd</sup> level  
can be combined

```
SELECT Name, maxima
FROM (SELECT *
      FROM Payroll P,
      (SELECT Job,
       MAX(Salary) AS maxima
       FROM Payroll
       GROUP BY Job) M
      WHERE P.Job = M.Job)
WHERE Salary = maxima
```




# The Witnessing Problem [More] Simplified

- Hold up, too complicated!
- 1<sup>st</sup> & 2<sup>nd</sup> level can be combined

```
SELECT Name, maxima
FROM Payroll P,
  (SELECT Job, MAX(Salary) AS maxima
   FROM Payroll
   GROUP BY Job) M
WHERE P.Job = M.Job AND Salary = maxima
```

```
SELECT Name, maxima
FROM (SELECT *
      FROM Payroll P,
        (SELECT Job,
          MAX(Salary) AS maxima
         FROM Payroll
         GROUP BY Job) M
      WHERE P.Job = M.Job)
WHERE Salary = maxima
```



- Ahh, that's better. Can we combine further?
- Well, yes, but it results in the self-join query we started with...
  - Need 2<sup>nd</sup> copy of Payroll
  - Outer WHERE becomes HAVING (for filter after aggregation)

# The Witnessing Problem [More] Simplified

## ■ Hold up, too complicated!

■ 1st & 2nd Level

Tradeoff between:

1. Complicated query w/o subqueries
2. "Simpler" query w/ subqueries

```
SELECT Name, maxima
FROM Payroll P,
     (SELECT Job, MAX(Salary) AS maxima
      FROM Payroll
      GROUP BY Job) M
WHERE P.Job = M.Job AND Salary = maxima
```

```
SELECT Name, maxima
FROM (SELECT *
      FROM Payroll P,
           (SELECT Job,
                MAX(Salary) AS maxima
            FROM Payroll
            GROUP BY Job) M
      WHERE P.Job = M.Job)
WHERE Salary = maxima
```


## ■ Ahh, that's better. Can we combine further?

- Well, yes, but it results in the self-join query we started with...
  - Need 2<sup>nd</sup> copy of Payroll
  - Outer WHERE becomes HAVING (for filter after aggregation)

# The Witnessing Problem Abstracted

- **Option #3: Save an intermediate result**
  - Good for abstraction
  - Good for reuse

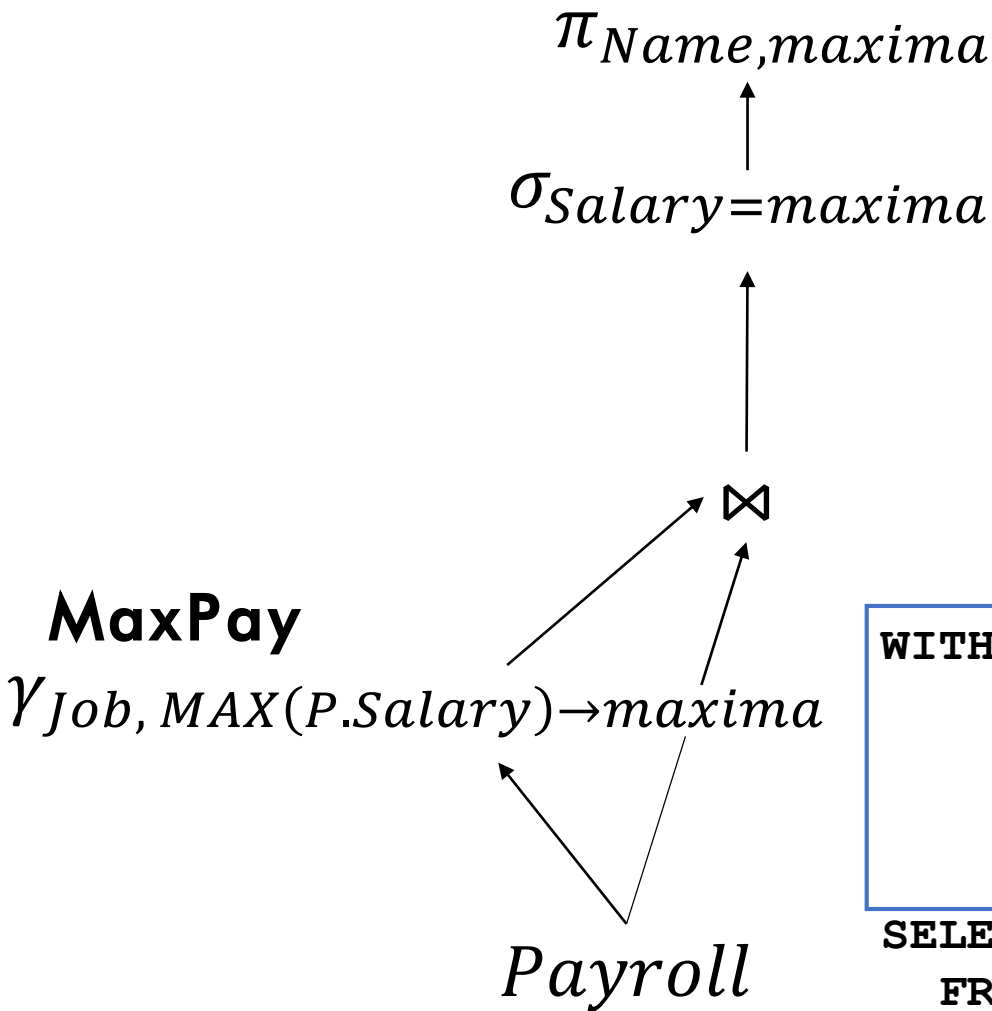
```
SELECT Name, maxima
FROM Payroll P,
  (SELECT Job,
    MAX(Salary) AS maxima
  FROM Payroll
  GROUP BY Job) M
WHERE P.Job = M.Job AND
      Salary = maxima
```



```
WITH MaxPay AS
  (SELECT Job AS Job,
    MAX(Salary) AS maxima
  FROM Payroll
  GROUP BY Job)
```

```
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.maxima
```

# The Witnessing Problem Abstracted



**MaxPay**

Job	Salary
TA	60000
Prof	100000

```
WITH MaxPay AS
    (SELECT Job AS Job,
            MAX(Salary) AS maxima
     FROM Payroll
     GROUP BY Job)
```

```
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.maxima
```

# The Witnessing Problem Summary

Option 1:  
Build larger  
query from  
subqueries

```
SELECT Name, maxima
FROM (SELECT *
      FROM Payroll P,
      (SELECT Job, MAX(Salary) AS maxima
       FROM Payroll
       GROUP BY Job) M
      WHERE P.Job = M.Job)
WHERE Salary = maxima
```

Option 2:  
Write a  
fancy big  
query

```
SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```

Option 3:  
Abstract a useful  
intermediate  
result

```
WITH MaxPay AS
  (SELECT Job AS Job, MAX(Salary) AS maxima
   FROM Payroll
   GROUP BY Job)
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND P.Salary = MP.maxima
```

# The Witnessing Problem Summary



```
SELECT Name, maxima
FROM (SELECT *
      FROM Payroll P,
      (SELECT Job, MAX(Salary) AS maxima
      FROM Payroll
      GROUP BY Job) M
      WHERE P.Job = M.Job)
WHERE Salary = maxima
```

```
SELECT P1.Name, MAX(P2.Salary)
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P2.Job, P1.Salary, P1.Name
HAVING P1.Salary = MAX(P2.Salary)
```

```
WITH MaxPay AS
      (SELECT Job AS Job, MAX(Salary) AS maxima
      FROM Payroll
      GROUP BY Job)
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND P.Salary = MP.maxima
```

# The Witnessing Problem Summary

psst... later we learn a 4<sup>th</sup> option:  
correlated subquery in the WHERE clause

So many options!

- With a perfect optimizer, all options are equivalent in speed; choose the easiest to write
- In practice, how you write a query affects speed

# General Subqueries

- Let's step back
- Subqueries can be interpreted as
  - **single values**
    - Single-tuple, single-attribute (1x1) relation
    - Returned as part of a tuple
  - **whole relations**
    - Used as input for another query
    - Checked for containment of a value



# Set Operations: SQL

- **SQL / RA mimics set theory**
  - Bag = duplicates allowed
  - **UNION (ALL)** → set union (bag union)
  - **INTERSECT (ALL)** → set intersection (bag intersection)
  - **EXCEPT (ALL)** → set difference (bag difference)
- **SQL Server Management Studio 2017**
  - INTERSECT ALL not supported
  - EXCEPT ALL not supported



# Set Operations: RA

## ■ SQL / RA mimics set theory

- Bag = duplicates allowed
- **UNION (ALL)** → set union (bag union)
- **INTERSECT (ALL)** → set intersection (bag intersection)
- **EXCEPT (ALL)** → set difference (bag difference)

Bag operators in RA are rare.

I've seen  $\cup$  for bag union.

2-arg Union op (Symbol "cup")  
Compute set union  
on matching schema

$A \cup B$

2-arg Intersection op (Symbol "cap")  
Compute set intersection  
on matching schema

$A \cap B$

2-arg Difference op (Backslash)  
Compute set difference  
on matching schema

$A \setminus B$

# Union

$\{ (123, \text{Jack}, \text{TA}, 50000), (345, \text{Allison}, \text{TA}, 60000), (567, \text{Magda}, \text{Prof}, 90000), (789, \text{Dan}, \text{Prof}, 100000) \}$   $\cup$   $\{ (987, \text{Gibbs}, \text{TA}, 60000), (423, \text{Shana}, \text{Prof}, 60000) \}$

$=$   $\{ (123, \text{Jack}, \text{TA}, 50000), (345, \text{Allison}, \text{TA}, 60000), (567, \text{Magda}, \text{Prof}, 90000), (789, \text{Dan}, \text{Prof}, 100000), (987, \text{Gibbs}, \text{TA}, 60000), (423, \text{Shana}, \text{Prof}, 60000) \}$

$\{ (123, \text{Jack}, \text{TA}, 50000), (345, \text{Allison}, \text{TA}, 60000), (567, \text{Magda}, \text{Prof}, 90000), (789, \text{Dan}, \text{Prof}, 100000) \}$   $\cup$   $\{ (987, \text{Gibbs}), (423, \text{Shana}) \}$

$= ???$

Schema mismatch!

# Set Operations

- SQL set-like operators slap two queries together (not really a subquery...)



# Subqueries in FROM

- Equivalent to a WITH subquery
- Use: Solve sub-problem to later join / evaluate

```
WITH MaxPay AS
    (SELECT P1.Job AS Job,
             MAX(P1.Salary) AS Salary
     FROM Payroll AS P1
     GROUP BY P1.Job)
SELECT P.Name, P.Salary
FROM Payroll AS P, MaxPay AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.Salary
```

Syntactic  
sugar

---

```
SELECT P.Name, P.Salary
FROM Payroll AS P, (SELECT P1.Job AS Job,
                           MAX(P1.Salary) AS Salary
                     FROM Payroll AS P1
                     GROUP BY P1.Job) AS MP
WHERE P.Job = MP.Job AND
      P.Salary = MP.Salary
```

# Subqueries in SELECT

- **Must return a single value**
  - single-tuple single-attribute

# Uncorrelated Subqueries in SELECT

```
SELECT Name,  
       (SELECT AVG(Salary)  
        FROM Payroll) AS AvgSal  
FROM Payroll
```

Uncorrelated =  
independent of outer query



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Uncorrelated Subqueries in SELECT

Name	AvgSal
Jack	75000
Allison	75000
Magda	75000
Dan	75000

- Hold up, we've seen this before!
- Transform to FROM clause



```
SELECT Name,  
       (SELECT AVG(Salary)  
        FROM Payroll) AS AvgSal  
FROM Payroll
```



UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



# Uncorrelated Subqueries in SELECT

Name	AvgSal
Jack	75000
Allison	75000
Magda	75000
Dan	75000

```
SELECT Name, AvgSal
FROM Payroll,
     (SELECT AVG(Salary)
      AS AvgSal)
FROM Payroll
```

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

$\pi_{Name, AvgSal}$

$\times$

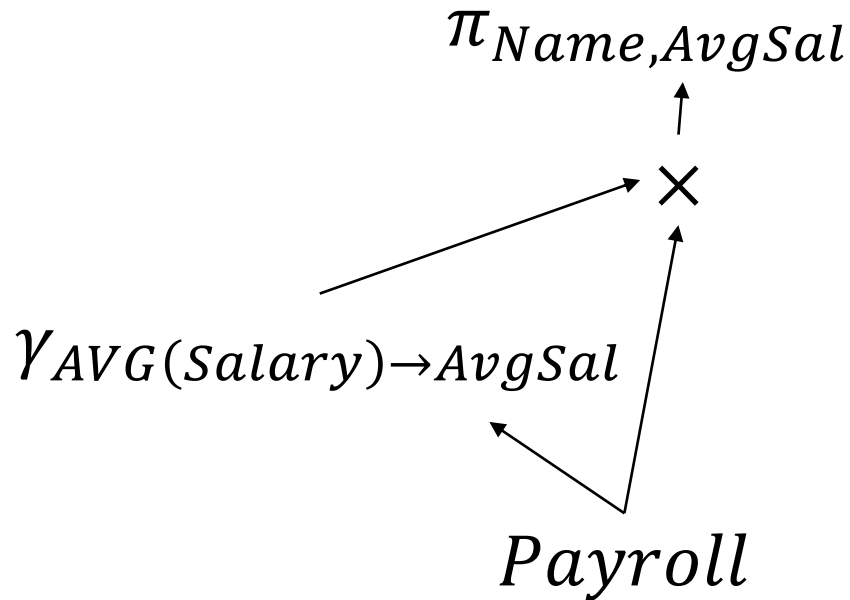
AvgSal

75000

$\gamma_{AVG(Salary) \rightarrow AvgSal}$

```
SELECT AVG(Salary)
      AS AvgSal
FROM Payroll
```

# Uncorrelated Subqueries in SELECT



- Uncorrelated SELECT subqueries don't add much
- Can convert to FROM subqueries

# Correlated Subqueries in SELECT

- Must return a single value
- Used to compute an associated value

For each person find the average salary of their job:

```
SELECT P.Name,  (SELECT AVG(P1.Salary)
                  FROM Payroll AS P1
                  WHERE P.Job = P1.Job)
                AS AvgSal
FROM Payroll AS P
```

# Correlated Subqueries in SELECT

- Must return a single value
- Used to compute an associated value

For each person find the average salary of their job:

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job)
                AS AvgSal
FROM Payroll AS P
```

**Correlated subquery!**  
The entire subquery is  
recomputed for each tuple

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

## Output

Name	AvgSal
------	--------

## Payroll P

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

## Output

Name	AvgSal
------	--------

## Payroll P

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

## Payroll P1

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

## Output

Name	AvgSal
------	--------

## Payroll P

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

## Payroll P1

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

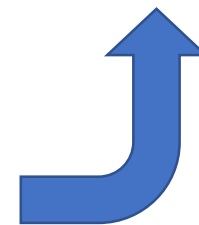
**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Output**

Name	AvgSal
Jack	55000

55000





# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



**Output**

Name	AvgSal
Jack	55000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Payroll P1**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Output**

Name	AvgSal
Jack	55000
Allison	55000

55000



# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



**Output**

Name	AvgSal
Jack	55000
Allison	55000
Magda	95000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



**Output**

Name	AvgSal
Jack	55000
Allison	55000
Magda	95000
Dan	95000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



**Output**

Name	AvgSal
Jack	55000
Allison	55000
Magda	95000
Dan	95000

# Correlated Subqueries in SELECT

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

**Output**

Name	AvgSal
Jack	55000
Allison	55000
Magda	95000
Dan	95000

- Hold up! How do we draw an RA Plan?
- There is no “evaluate a different subquery for each tuple” operator...
- Solution: de-correlate by adding a join



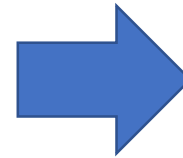
## Design an equivalent query without a SELECT subquery

For each person find the average salary of their job:

```
SELECT P.Name, (SELECT AVG(P1.Salary)
                FROM Payroll AS P1
                WHERE P.Job = P1.Job) AS AvgSal
FROM Payroll AS P
```

**Payroll P**

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000



**Output**

Name	AvgSal
Jack	55000
Allison	55000
Magda	95000
Dan	95000

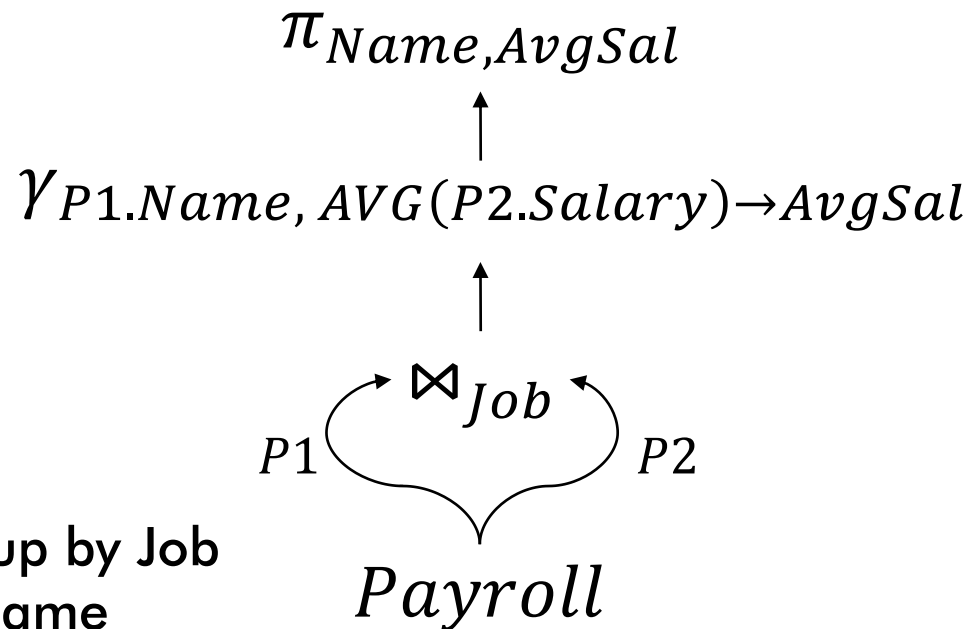
# Solution #1

For each person find the average salary of their job

```
SELECT P1.Name, AVG(P2.Salary) AS AvgSal
FROM Payroll AS P1, Payroll AS P2
WHERE P1.Job = P2.Job
GROUP BY P1.Name
```

Idea:

1. Self-join Payroll
2. Use one copy to aggregate, group by Job
3. Use one copy to keep original Name





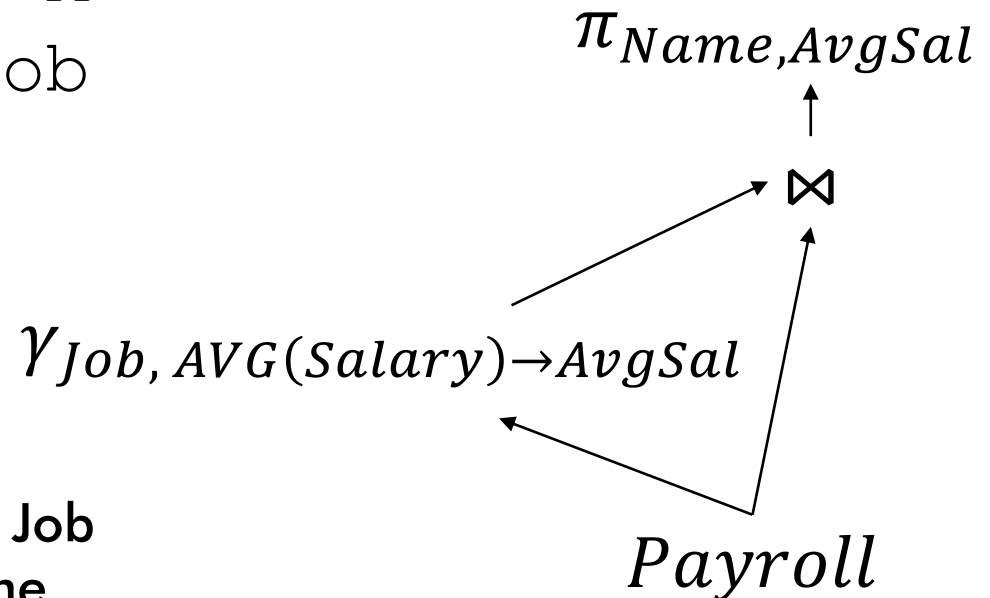
# Solution #2

For each person find the average salary of their job

```
SELECT P.Name, M.AvgSal
FROM Payroll P,
      (SELECT Job, AVG(Salary) AS AvgSal
       FROM Payroll
       GROUP BY Job) M
WHERE P.Job = M.Job
```

Idea:

1. Compute aggregate, group by Job
2. Join to original Payroll for Name



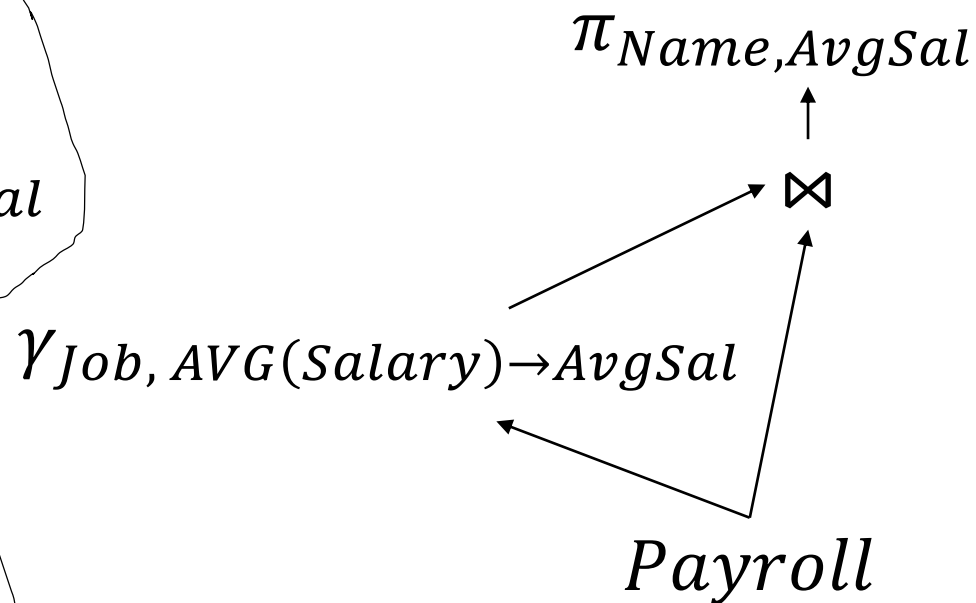
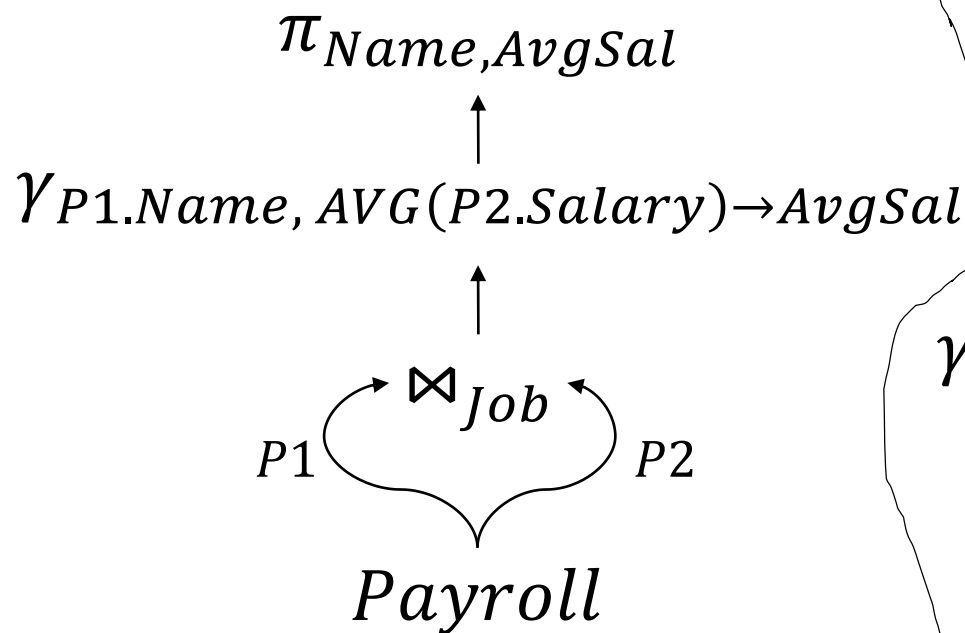
# Preview: RA Equivalence

For each person find the average salary of their job

```
SELECT P1.Name,  
       AVG(P2.Salary) AS AvgSal  
FROM Payroll AS P1, Payroll AS P2  
WHERE P1.Job = P2.Job  
GROUP BY P1.Name
```

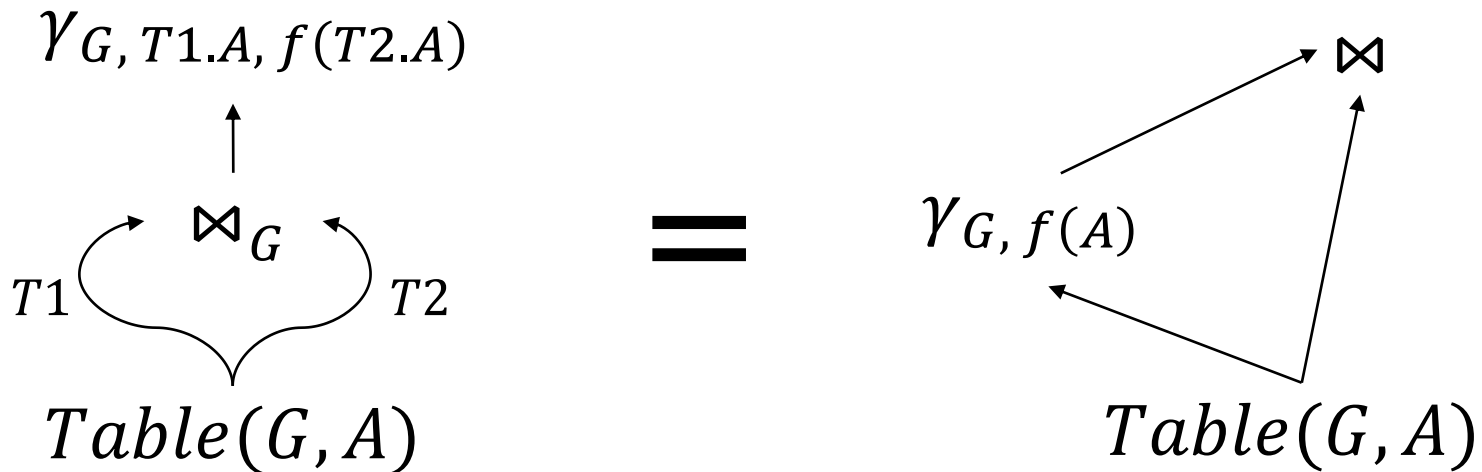
```
SELECT P.Name, M.AvgSal  
FROM Payroll P,  
     (SELECT Job, AVG(Salary) AS AvgSal  
      FROM Payroll  
      GROUP BY Job) M  
WHERE P.Job = M.Job
```

=



# Preview: RA Equivalence

- General form: “Push aggregate into join” rewrite
- Speedup
- Proof using formal RA



# Subqueries in SELECT: Special Cases

For each person find the number of cars they drive

```
SELECT P.Name,  
       (SELECT COUNT(R.Car)  
        FROM Regist AS R  
        WHERE P.UserID = R.UserID)  
FROM Payroll AS P
```

```
SELECT P.Name, COUNT(R.Car)  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID  
GROUP BY P.Name
```

= ?

Same? Discuss!

(RA impossible with  
correlated SELECT subquery)

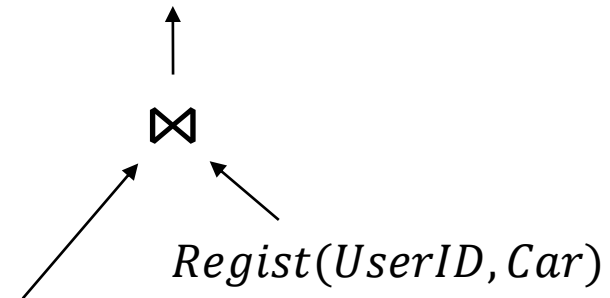
Payroll

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

Regist

UserID	Car
123	Charger
567	Civic
567	Pinto

$\gamma_{Name, Count(Car)}$



*Payroll*(UserID, Name, Job, Salary)

# Subqueries in SELECT: Special Cases

For each person find the number of cars they drive

```
SELECT P.Name,  
       (SELECT COUNT(R.Car)  
        FROM Regist AS R  
        WHERE P.UserID = R.UserID)  
FROM Payroll AS P
```

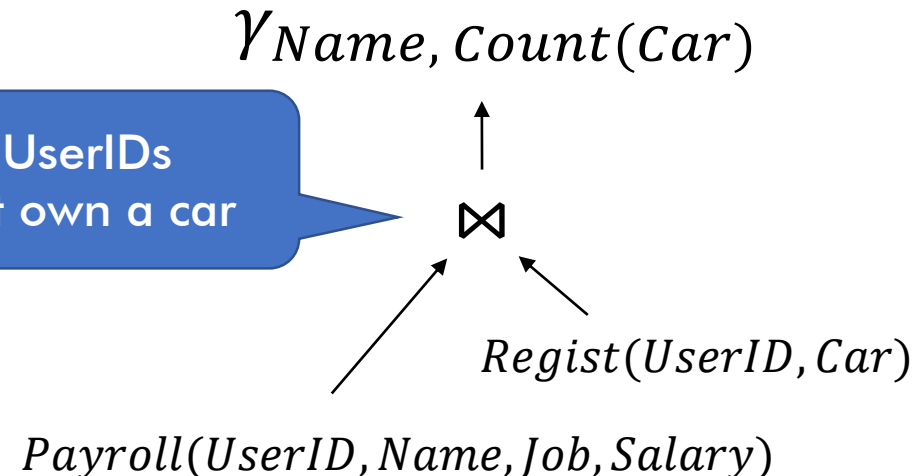
Returns 0 for UserIDs  
that don't own a car

≠

```
SELECT P.Name, COUNT(R.Car)  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID  
GROUP BY P.Name
```

User OUTER JOIN to  
retain UserIDs  
that don't own a car

Drops UserIDs  
that don't own a car



# Subqueries in SELECT: Special Cases

For each person find the number of cars they drive

```
SELECT P.Name,  
       (SELECT COUNT(R.Car)  
        FROM Regist AS R  
        WHERE P.UserID = R.UserID)  
FROM Payroll AS P
```

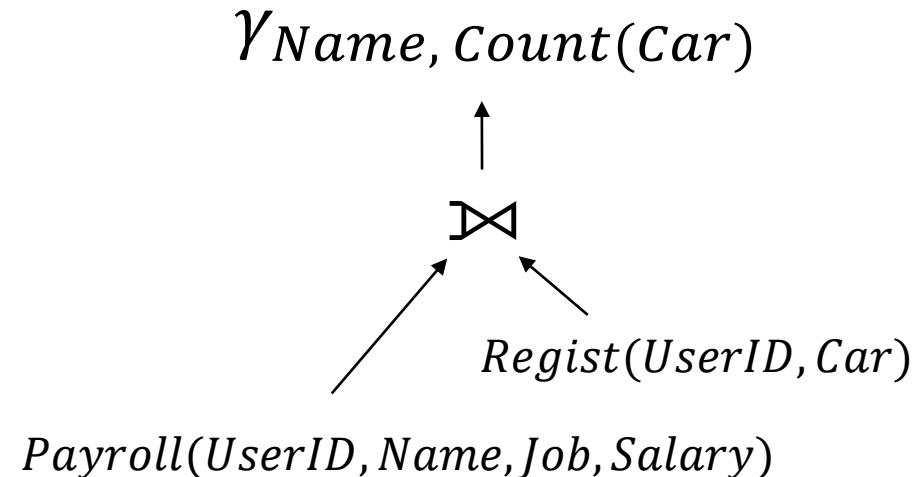
```
SELECT P.Name, COUNT(R.Car)  
FROM Payroll AS P  
LEFT OUTER JOIN Regist AS R  
ON P.UserID = R.UserID  
GROUP BY P.Name
```

UserID	COUNT(Car)
123	1
345	0
567	2
789	0

=

UserID	Name	Job	Salary
123	Jack	TA	50000
345	Allison	TA	60000
567	Magda	Prof	90000
789	Dan	Prof	100000

UserID	Car
123	Charger
567	Civic
567	Pinto



# Subqueries in WHERE/HAVING

- **Basic use: compare to scalar** (single tuple single attribute)

Find the name who earns the highest salary

```
SELECT P.Name
  FROM Payroll P
 WHERE P.Salary = (SELECT MAX(P1.Salary)
                  FROM Payroll P1)
```

General strategy:  
Join the condition into  
a new column

- **Uncorrelated subquery**
- **Need to separate before can draw RA**
- **We've seen this before: simple witness problem**

# Subqueries in WHERE/HAVING

- **Basic use: compare to scalar** (single tuple single attribute)

Find the name who earns the highest salary  
*for each job type*

```
SELECT P.Name
FROM Payroll P
WHERE P.Salary = (SELECT MAX(P1.Salary)
                  FROM Payroll P1
                  WHERE P.Job = P1.Job)
```

- **Correlated subquery**
  - Like SELECT subquery, evaluated per tuple
- Like the witness problem



# Subqueries in WHERE/HAVING

## ■ Advanced keywords:

- ANY  $\rightarrow \exists$
  - ALL  $\rightarrow \forall$
  - (NOT) IN  $\rightarrow (\notin) \in$
  - (NOT) EXISTS  $\rightarrow$
- Use with a condition ( $=$ ,  $>$ ,  $<$ , ...)
- Check if a tuple is (not) part of a relation
- Check if a relation is (not) empty

# Subqueries in WHERE/HAVING

## ■ Advanced keywords:

SQLite does not support ANY or ALL

- ANY  $\rightarrow \exists$
  - ALL  $\rightarrow \forall$
  - (NOT) IN  $\rightarrow (\notin) \in$
  - (NOT) EXISTS  $\rightarrow$
- Use with a condition ( $=$ ,  $>$ ,  $<$ , ...)
- Check if a tuple is (not) part of a relation
- Check if a relation is (not) empty

Again: Find the name who earns the highest salary  
*for each job type*

```
SELECT P.Name
FROM Payroll AS P
WHERE P.Salary >= ALL (SELECT Salary
                        FROM Payroll
                        WHERE P.Job = Job)
```

# Subqueries in WHERE/HAVING

## ■ Advanced keywords:

- ANY  $\rightarrow \exists$
  - ALL  $\rightarrow \forall$
  - (NOT) IN  $\rightarrow (\notin) \in$
  - (NOT) EXISTS  $\rightarrow$
- Use with a condition ( $=$ ,  $>$ ,  $<$ , ...)
- Check if a tuple is (not) part of a relation
- Check if a relation is (not) empty

Find the name and salary of people who do not drive cars

# Subqueries in WHERE/HAVING

## ■ Advanced keywords:

- ANY  $\rightarrow \exists$
  - ALL  $\rightarrow \forall$
  - (NOT) IN  $\rightarrow (\notin) \in$
  - (NOT) EXISTS  $\rightarrow$
- Use with a condition ( $=$ ,  $>$ ,  $<$ , ...)
- Check if a tuple is (not) part of a relation
- Check if a relation is (not) empty

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT UserID
                        FROM Regist)
```

De-correlated! But not enough to draw RA (no "NOT IN" operator)

# Subqueries in WHERE/HAVING

## ■ Advanced keywords:

- ANY  $\rightarrow \exists$
  - ALL  $\rightarrow \forall$
  - (NOT) IN  $\rightarrow (\notin) \in$
  - (NOT) EXISTS  $\rightarrow$
- Use with a condition ( $=$ ,  $>$ ,  $<$ , ...)
- Check if a tuple is (not) part of a relation
- Check if a relation is (not) empty

Find the name and salary of people who do not drive cars

To convert to an RA Plan,  
rewrite using UNION, INTERSECT, or EXCEPT

```
WHERE P.UserID NOT IN (SELECT UserID  
                        FROM Regist)
```

De-correlated! But not enough to draw RA (no "NOT IN" operator)

# Subqueries in WHERE/HAVING

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary  
  FROM Payroll AS P  
 WHERE P.UserID NOT IN (SELECT UserID FROM Regist)
```

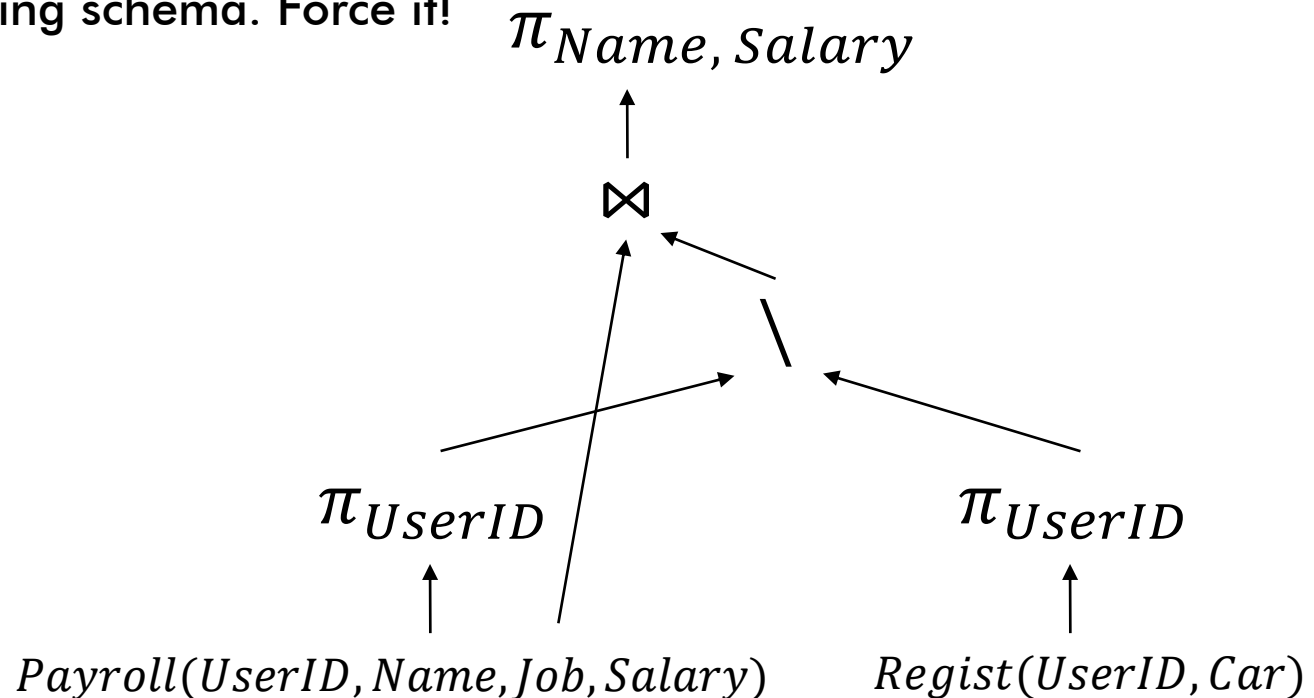
- This query takes Payroll UserIDs and subtracts Regist UserIDs
- Subtraction = EXCEPT
- EXCEPT requires matching schema. Force it!

# Subqueries in WHERE/HAVING

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT UserID FROM Regist)
```

- This query takes Payroll UserIDs and subtracts Regist UserIDs
- Subtraction = EXCEPT
- EXCEPT requires matching schema. Force it!



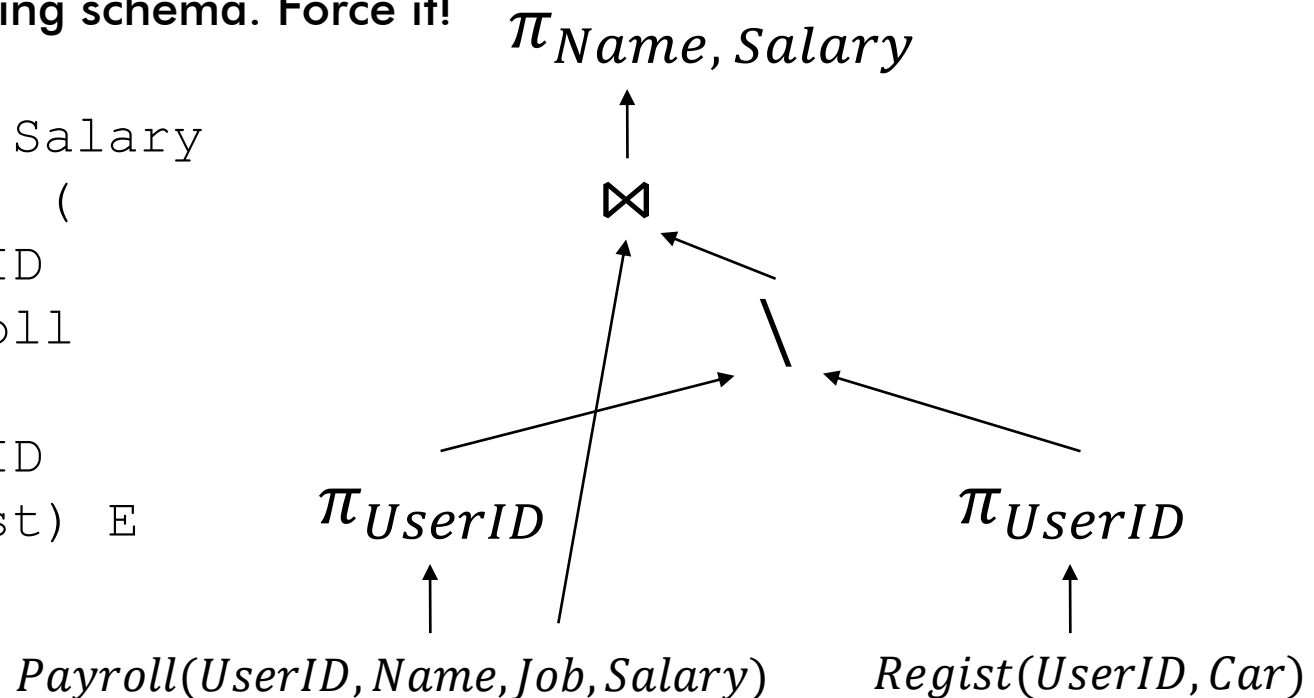
# Subqueries in WHERE/HAVING

Find the name and salary of people who do not drive cars

```
SELECT P.Name, P.Salary
FROM Payroll AS P
WHERE P.UserID NOT IN (SELECT UserID FROM Regist)
```

- This query takes Payroll UserIDs and subtracts Regist UserIDs
- Subtraction = EXCEPT
- EXCEPT requires matching schema. Force it!

```
SELECT P.Name, P.Salary
FROM Payroll P, (
    SELECT UserID
    FROM Payroll
EXCEPT
    SELECT UserID
    FROM Regist) E
WHERE P.UserID
    = E.UserID
```

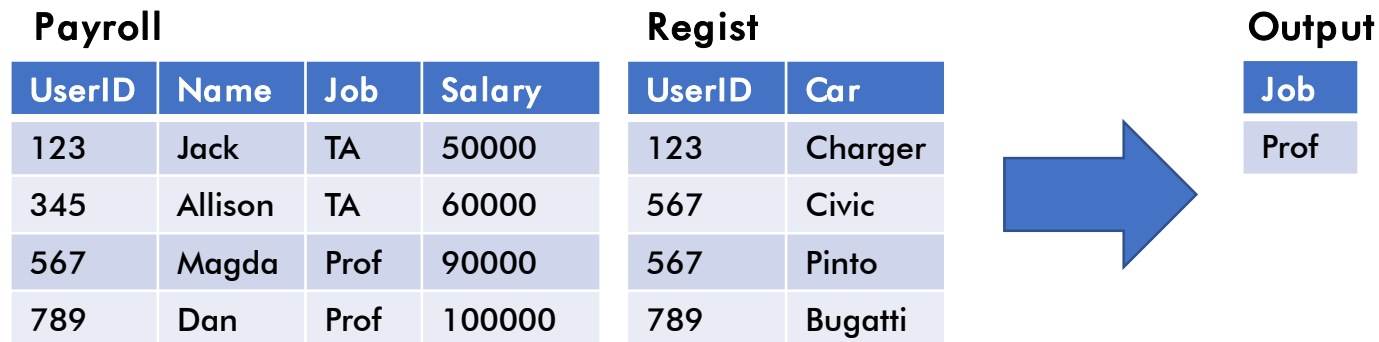




# Hard Cases: Universal Quantifiers

Find jobs whose employees **all** own cars

- All = “every employee must own a car”
- Hard to compute directly
- Try computing the negation!



# Hard Cases: Universal Quantifiers

Find jobs whose employees **all** own cars

→  $\forall$  employee  $e \in \text{job}$ ,  $e$  owns a car

$\forall$  = "for all"

Try computing the negation!

→  $\neg(\forall \text{ employee } e \in \text{job}, e \text{ owns a car})$

$\neg$  = "not"

→  $\exists \text{ employee } e \in \text{job}, \neg(e \text{ owns a car})$

→  $\exists \text{ employee } e \in \text{job}, e \text{ doesn't own a car}$

Find jobs with **an** employee who **doesn't** own a car

$\exists$  = "there exists (at least one)"

De Morgan's Law

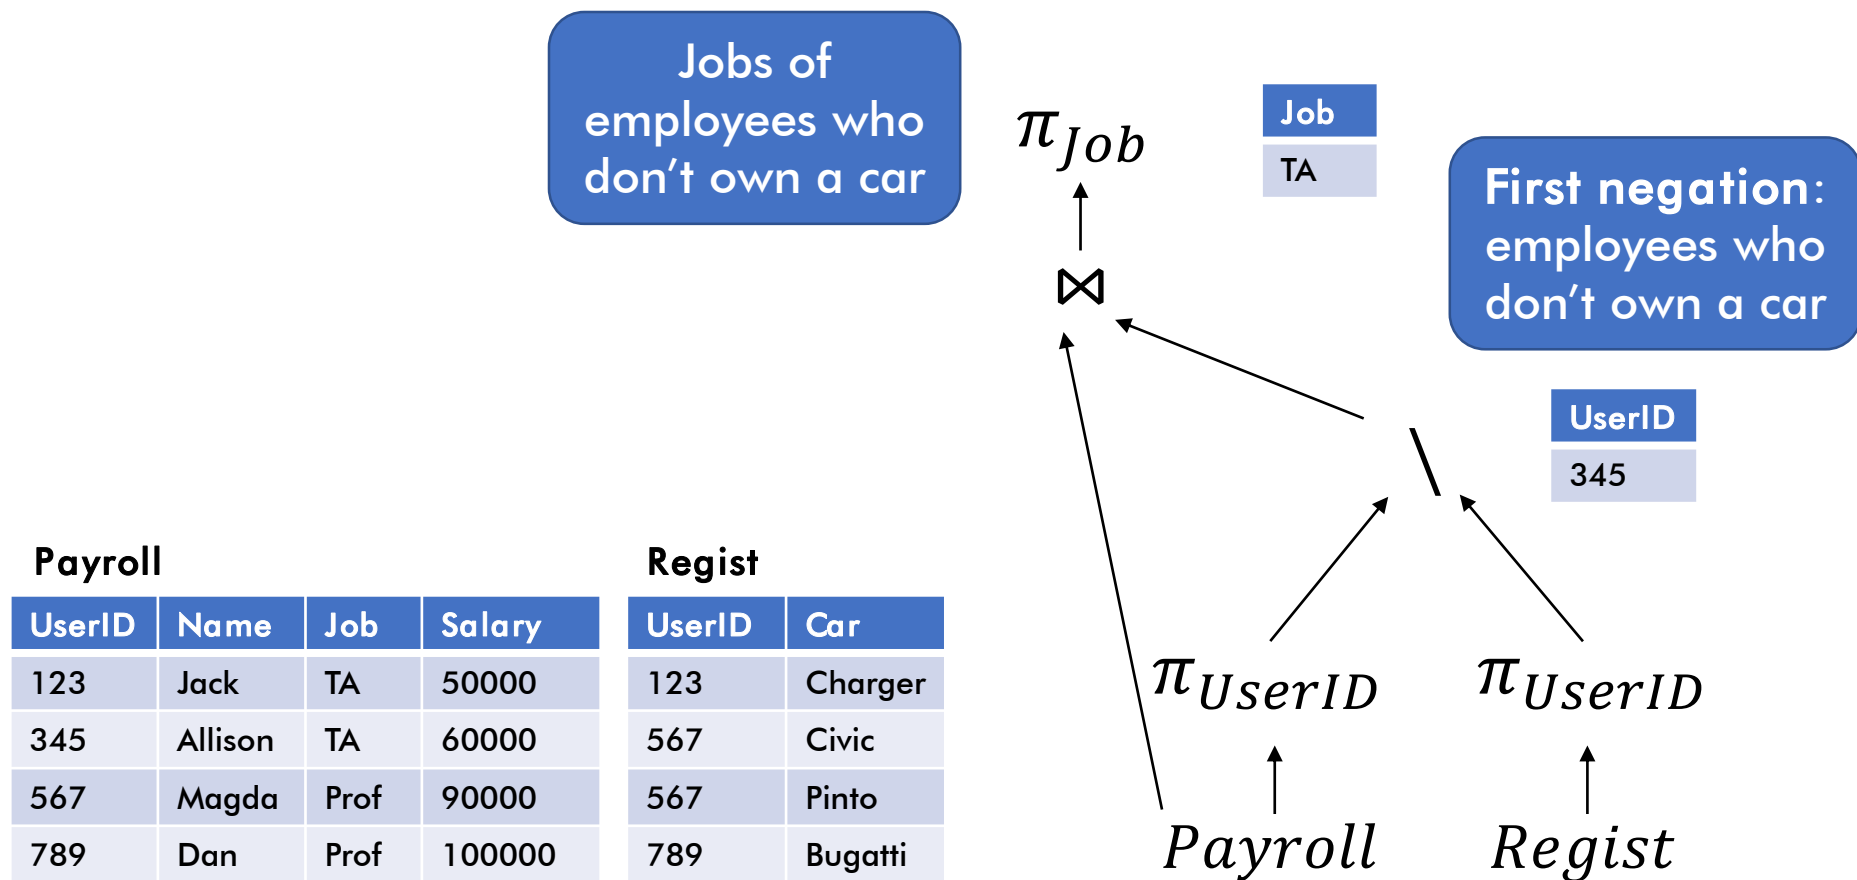
$$\neg(\forall x, c) = \exists x, \neg c$$

It's okay if you haven't seen logic symbols  $\forall$ ,  $\exists$ ,  $\neg$  before. Think about how to logically negate an English sentence

# Hard Cases: Universal Quantifiers

Find jobs whose employees *all* own cars

→ Find jobs with **an** employee who **doesn't** own a car



# Hard Cases: Universal Quantifiers

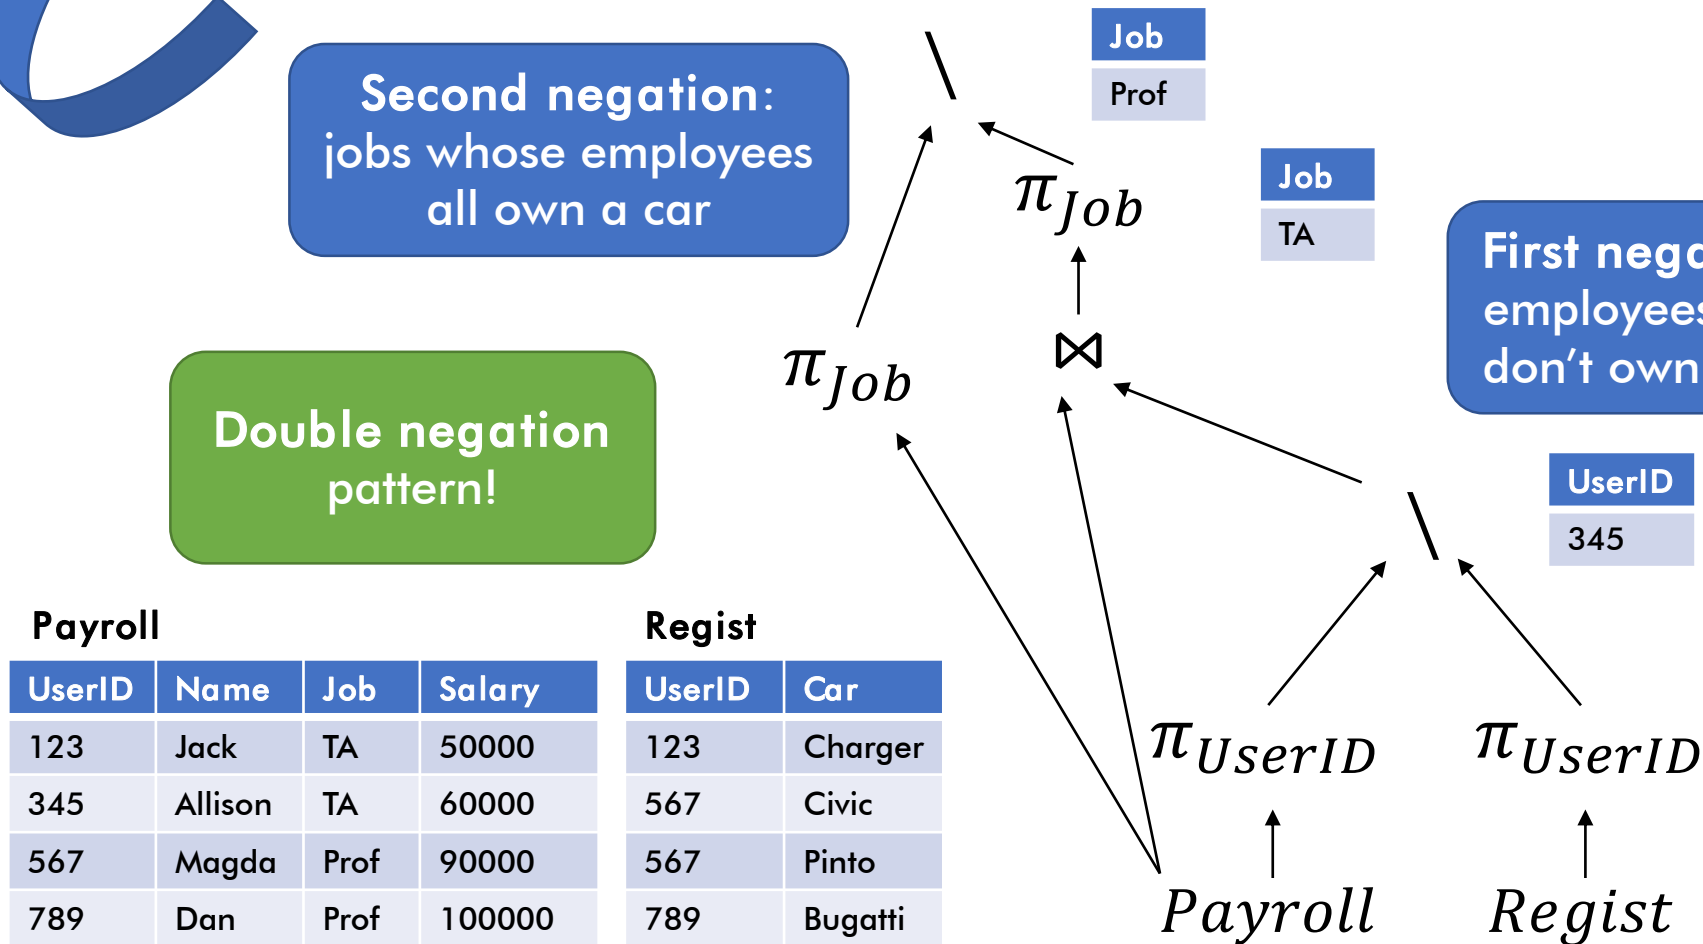
Find jobs whose employees *all* own cars

Find jobs with **an** employee who **doesn't** own a car

Second negation:  
jobs whose employees  
all own a car

Double negation  
pattern!

First negation:  
employees who  
don't own a car



# Hard Cases: Universal Quantifiers

- Watch out for universal quantifiers
  - Require more complex answer
- *Double negation* pattern often works

$$\forall = \neg \exists \neg$$